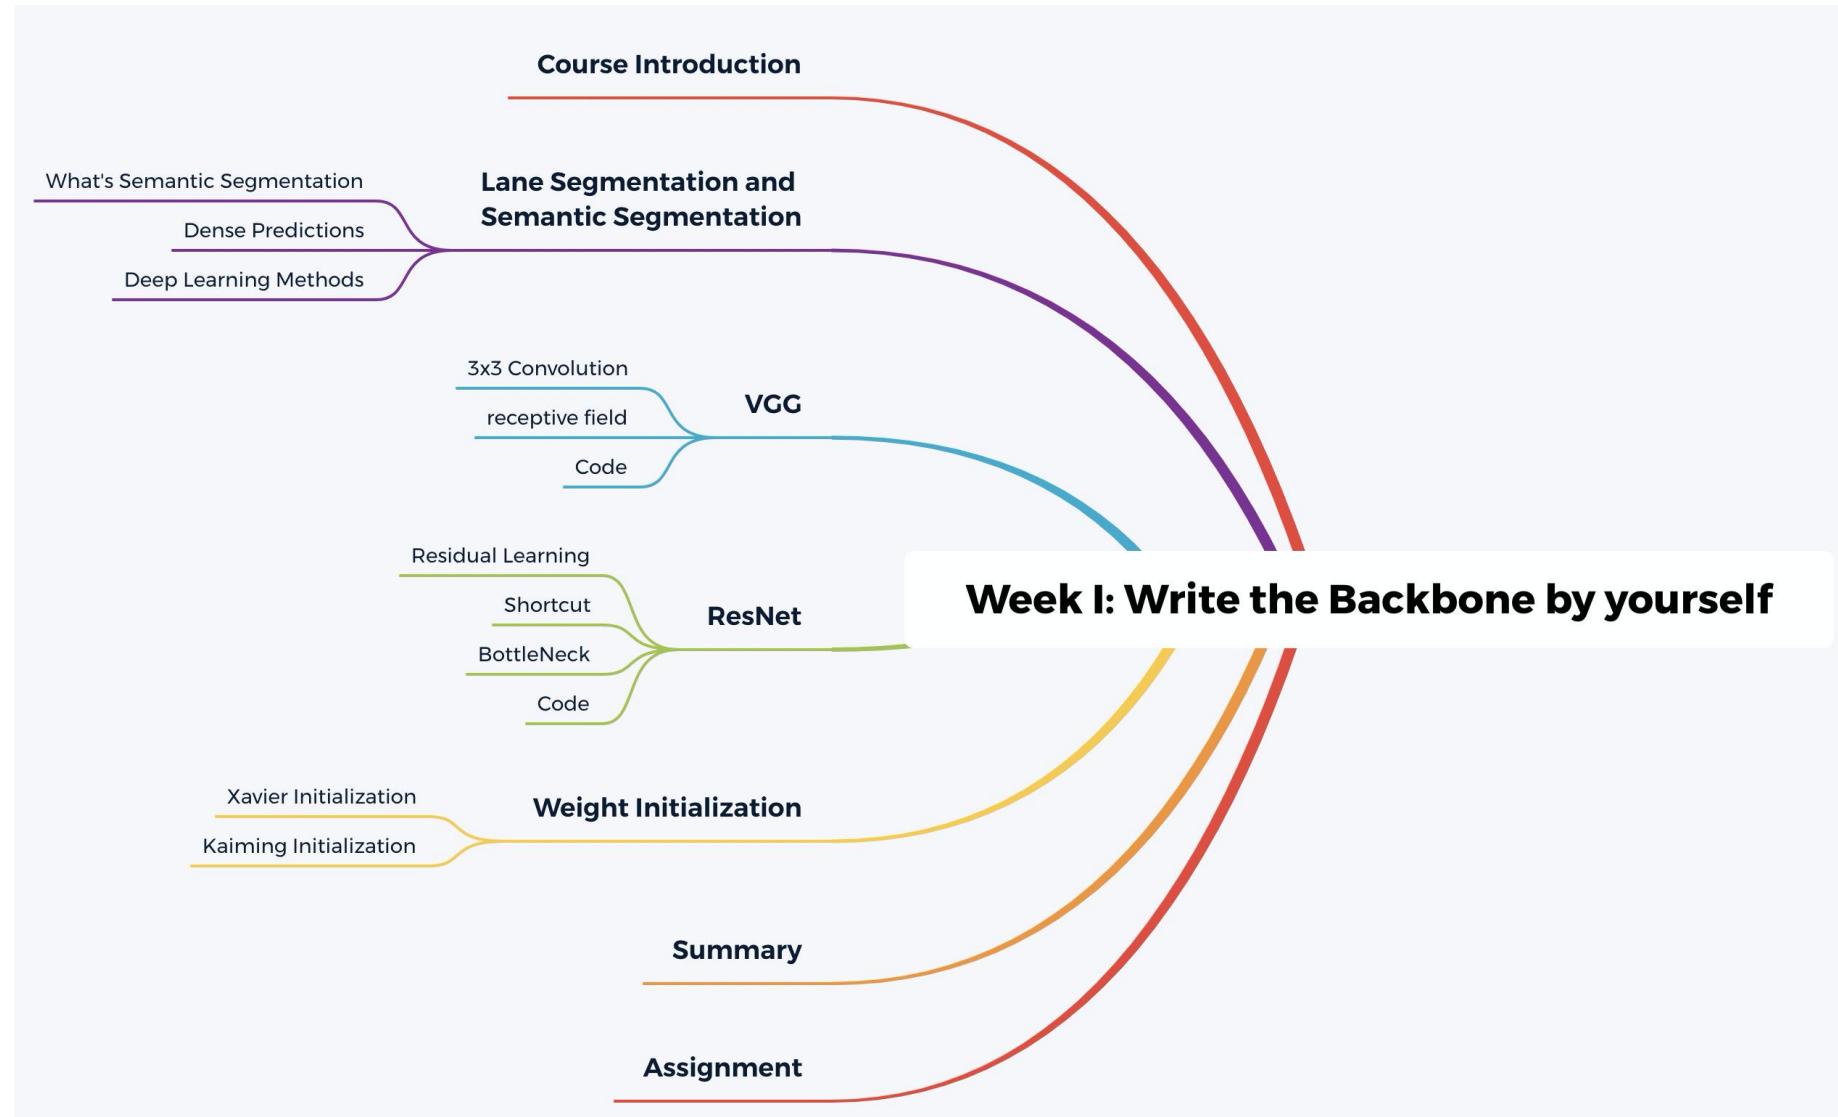


Lane Segmentation Week 1

HCT CV Class

主要内容



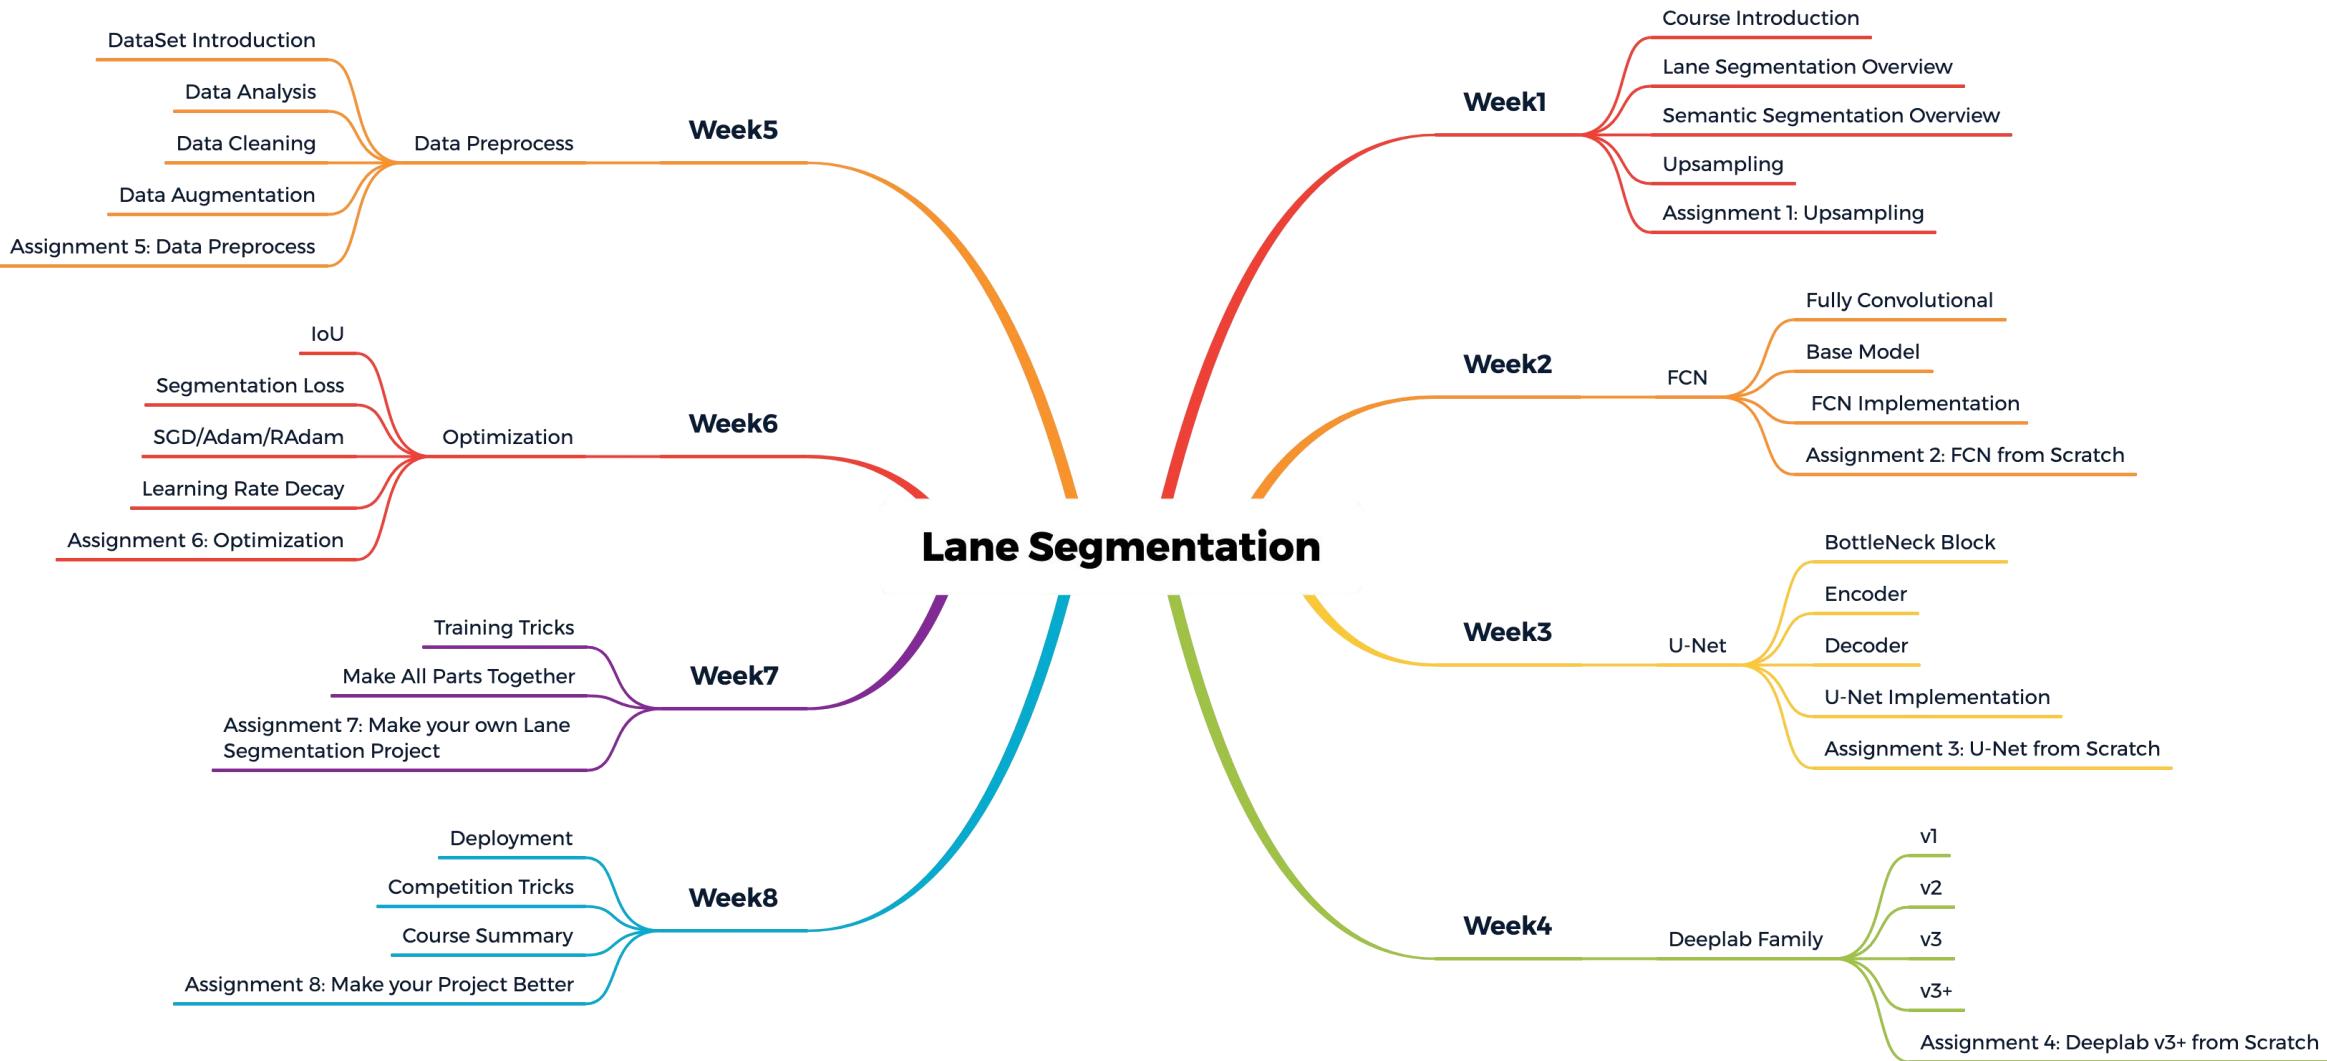
学习目标

- 了解Semantic Segmentation
- 掌握VGG
- 掌握ResNet
- 掌握Weight Initialization

课程介绍

- 详细讲解Lane Segmentation的原理和代码实现。通过课程学习可以掌握Semantic Segmentation的原理和经典模型的构建，掌握百度无人驾驶车道线数据集的预处理方法和优化训练技巧，不仅可以达到百度车道线挑战赛一流成绩，而且可以广泛应用于各种图像语义分割任务，对图像识别和目标检测等其他深度学习计算机视觉任务的水平也将会有很大提高。

Syllabus



预备知识

- Python
- Machine Learning
- CNN and Image Classification
- Tensorflow or Pytorch or PaddlePaddle...
- English
- Math

Project



无人车车道线检测挑战赛

[结束](#)

本次PaddlePaddle-无人车车道线检测挑战赛旨在为参赛者提供一定数量的准确的车道线标注数据，让更多的研究者参与并设计出高效、准确的检测算法，以此来共同推动无人车的发展，从而造福整个社会。

[已结束](#)

奖池：¥ 99000

报名人数: 743

[比赛介绍](#)[赛题说明](#)[提交结果](#)[我的团队](#)[排行榜](#)[讨论区](#)

竞赛简介

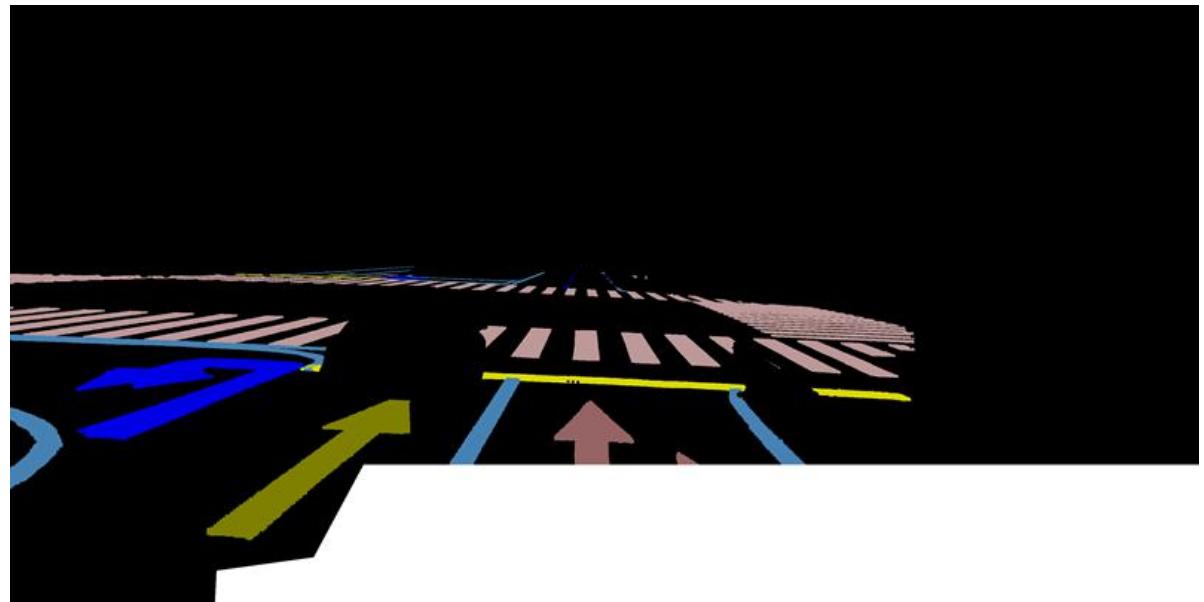
作为下一个改变人类生活方式的技术蓝海，无人驾驶车辆研究日益受到学术界和产业界的高度重视，其技术涉及认知科学、人工智能、控制科学、机械工程等交叉学科，是各种新兴技术的最佳验证平台，也是未来汽车发展的必然趋势。在无人驾驶汽车研究领域，带有车道线属性的高精地图是商业无人驾驶的一个非常关键的环节。截止目前，大多数的高精地图都是靠人工标注来完成的。人工标注不仅效率低，而且成本高不适宜大规模的商业应用。在高精地图的生产过程中，如何从图像中分割出不同类型的车道线是非常重要的一个环节。同时，准确分割车道线也会为将来高精地图的更新提供帮助。本次无人车车道线检测挑战赛旨在为参赛者提供一定数量的准确的车道线标注数据，让更多的研究者参与并设计出高效、准确的检测算法，以此来共同推动无人车的发展，从而造福整个社会。

赛程赛制

竞赛分为初赛、复赛两个阶段，各阶段参赛队伍必须按照要求按时、合规地提交参赛作品。

<https://aistudio.baidu.com/aistudio/competition/detail/5>

Lane Segmentation



<https://aistudio.baidu.com/aistudio/competition/detail/5>

CV Tasks

Classification



CAT

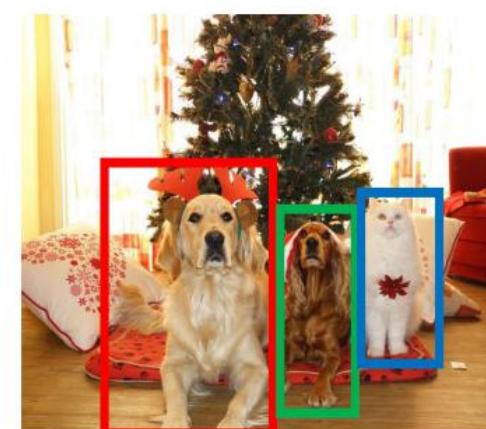
Semantic Segmentation



GRASS, CAT,
TREE, SKY

No spatial extent

Object Detection



DOG, DOG, CAT

Instance Segmentation



DOG, DOG, CAT

Multiple Object

This image is CC0 public domain

Semantic Segmentation



No spatial extent

Semantic Segmentation



GRASS, CAT,
TREE, SKY

No objects, just pixels

Object Detection



Multiple Object

Instance Segmentation



Semantic Segmentation

- Label each pixel in the image with a category label
- Don't differentiate instances, only care about pixels
- We can think of semantic segmentation as image classification at a pixel level

Semantic Labels



segmented →

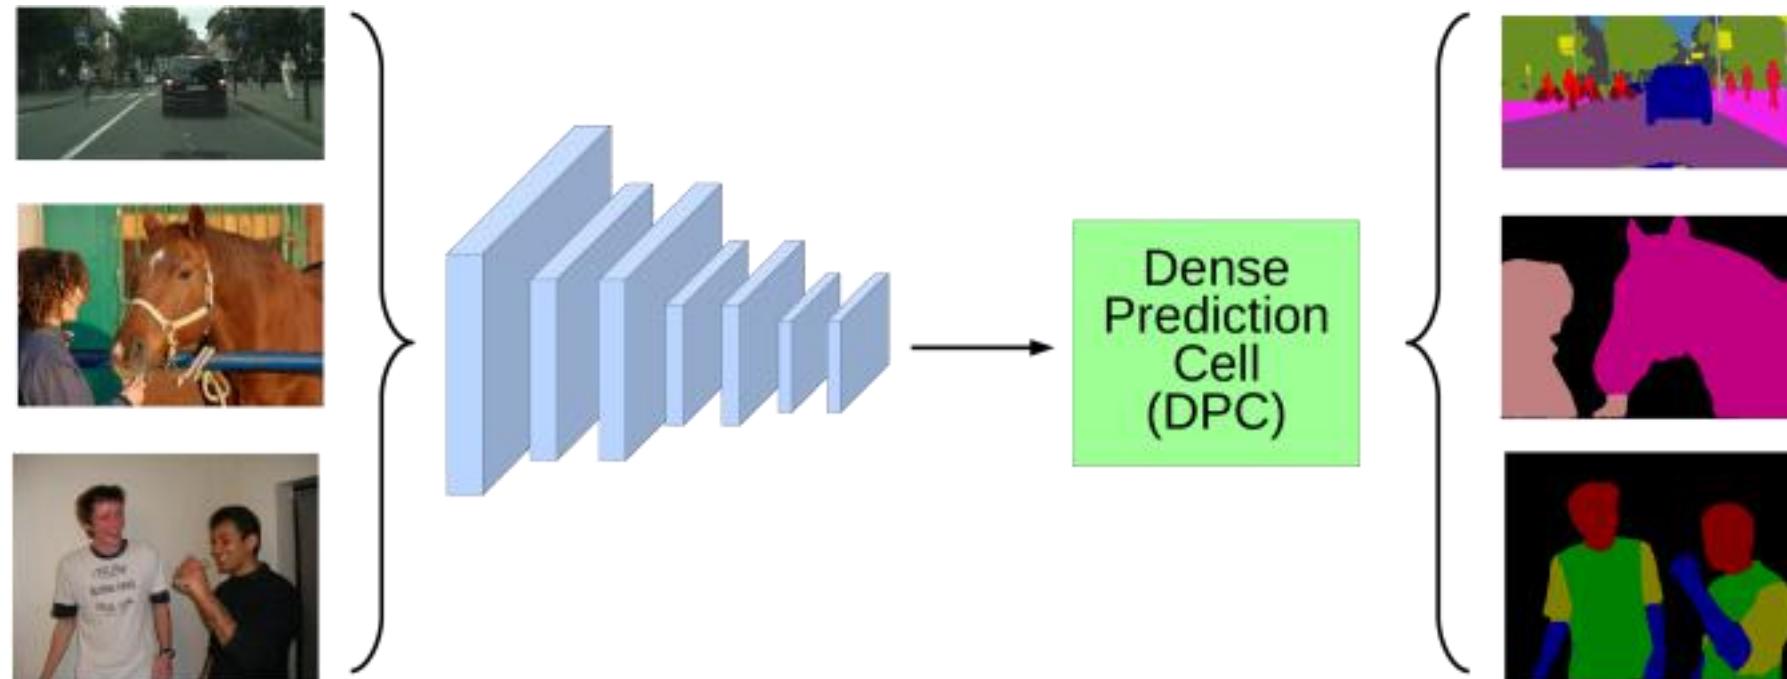
- 1: Person
- 2: Purse
- 3: Plants/Grass
- 4: Sidewalk
- 5: Building/Structures

Input

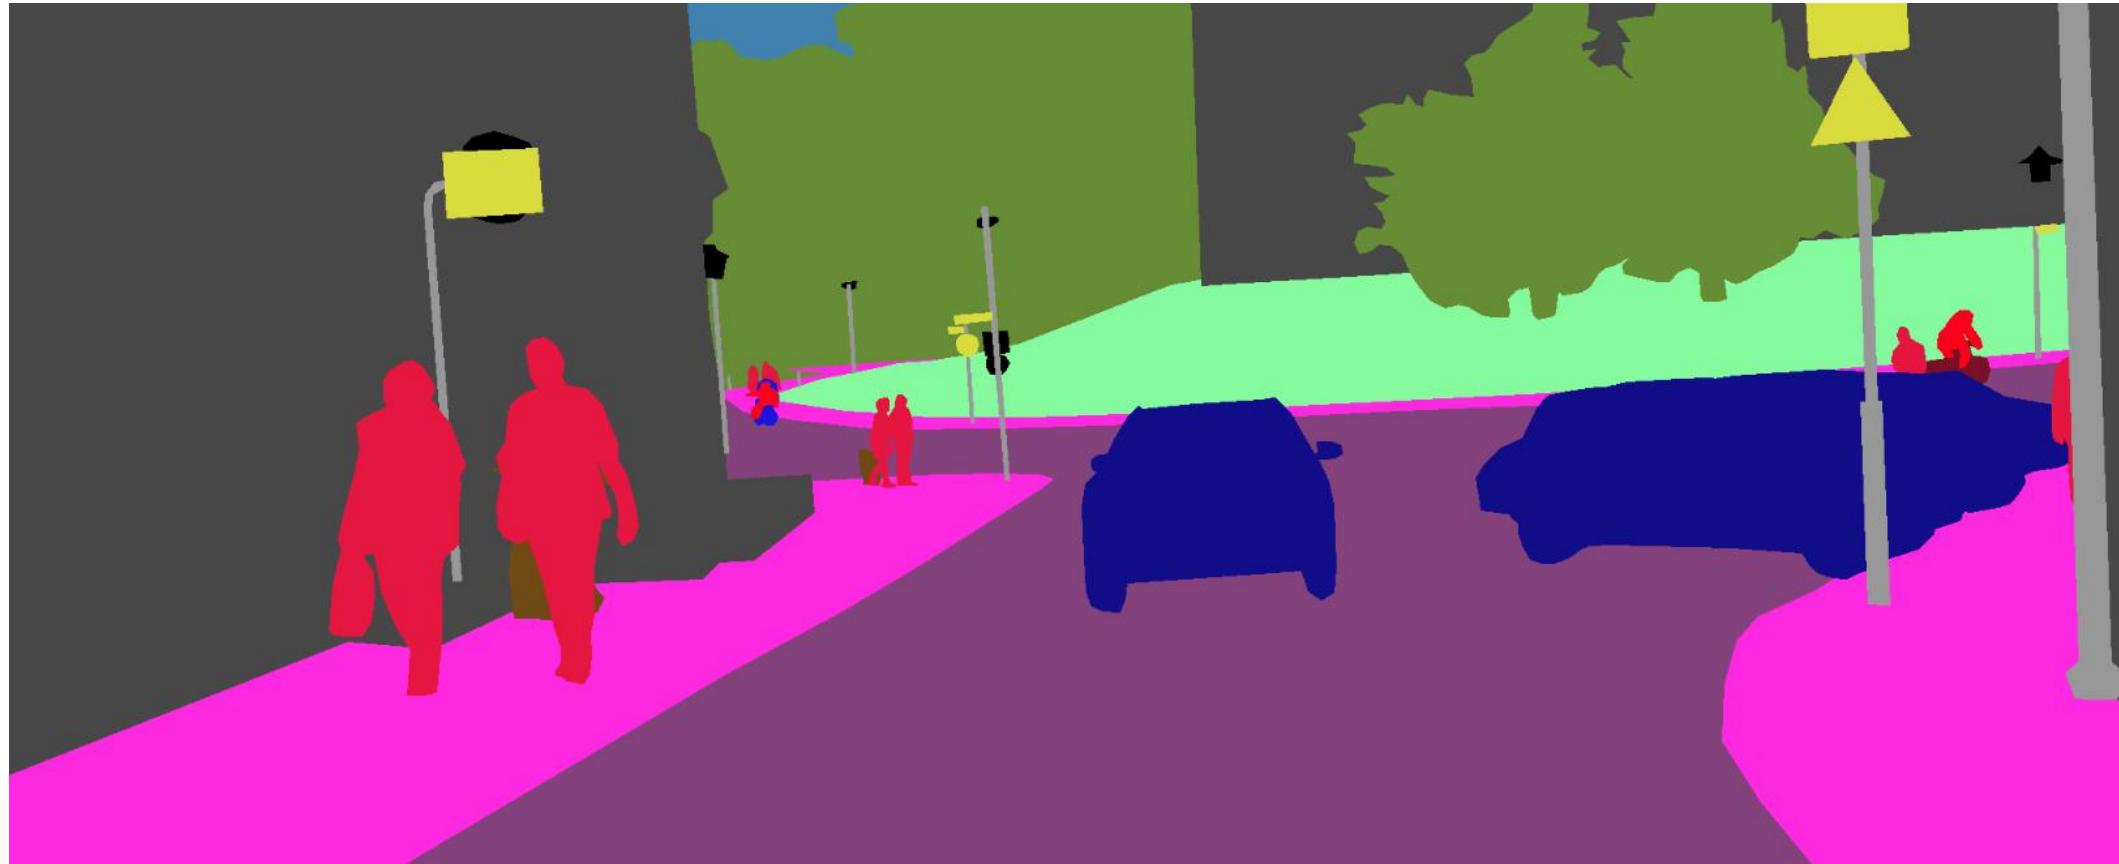
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	5	5	5	5	5
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	5	5	5	5	5
3	3	3	3	3	3	3	3	3	1	1	3	3	3	3	5	5	5	5	5	5
3	3	3	3	3	3	3	1	1	1	1	3	3	3	3	5	5	5	5	5	5
3	3	3	3	3	3	3	3	1	1	1	3	3	3	3	5	5	5	5	5	5
5	5	3	3	3	3	3	3	1	1	3	3	3	5	5	5	5	5	5	5	5
4	4	3	4	1	1	1	1	1	1	1	4	4	4	5	5	5	5	5	5	5
4	4	3	4	1	1	1	1	1	1	1	4	4	4	4	4	5	5	5	5	5
4	4	4	1	1	1	1	1	1	1	1	4	4	4	4	4	4	4	4	4	4
3	3	3	1	1	1	1	1	1	1	1	4	4	4	4	4	4	4	4	4	4
3	3	3	1	2	2	1	1	1	1	1	4	4	4	4	4	4	4	4	4	4
3	3	3	1	2	2	1	1	1	1	1	4	4	4	4	4	4	4	4	4	4

Semantic Labels

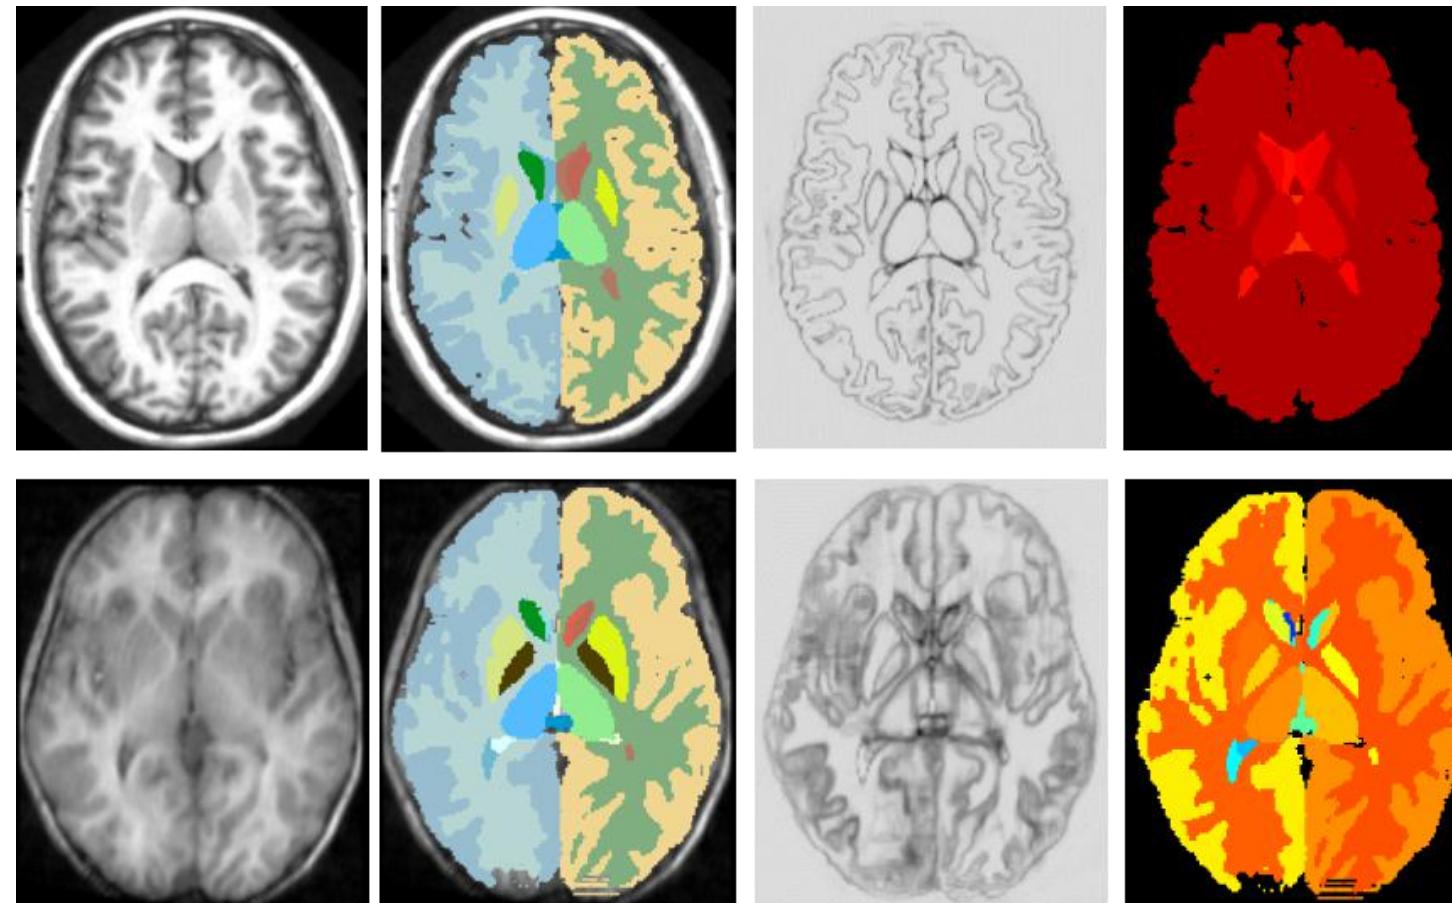
Dense Predictions



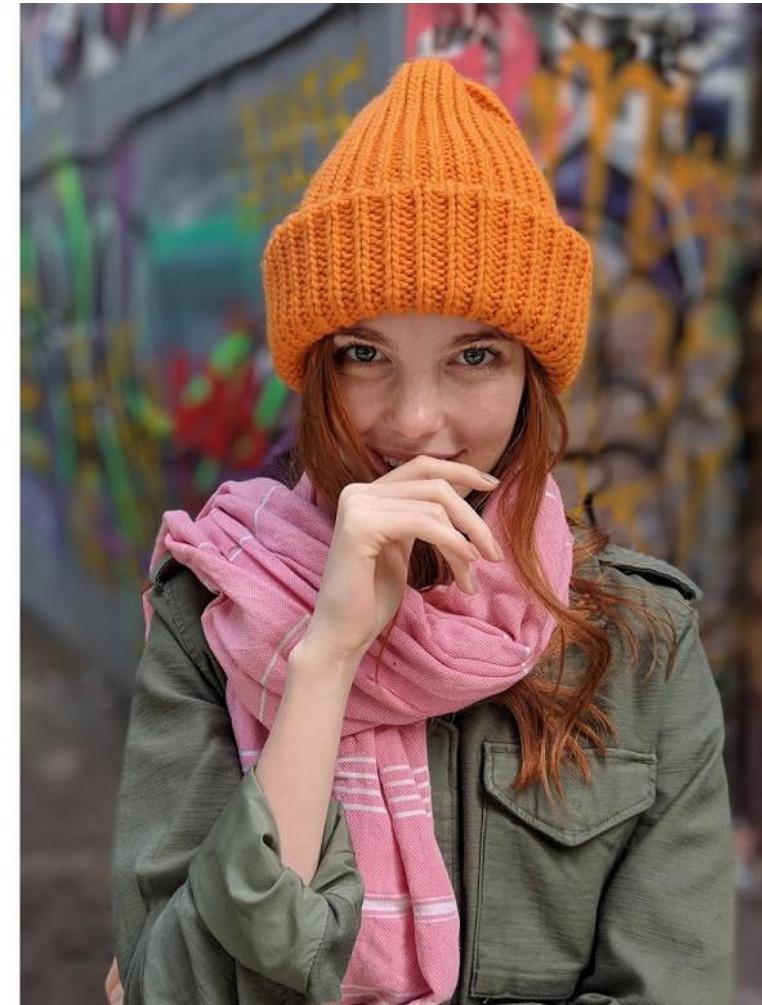
Auto Driving



medical image segmentation



shallow depth-of-field image



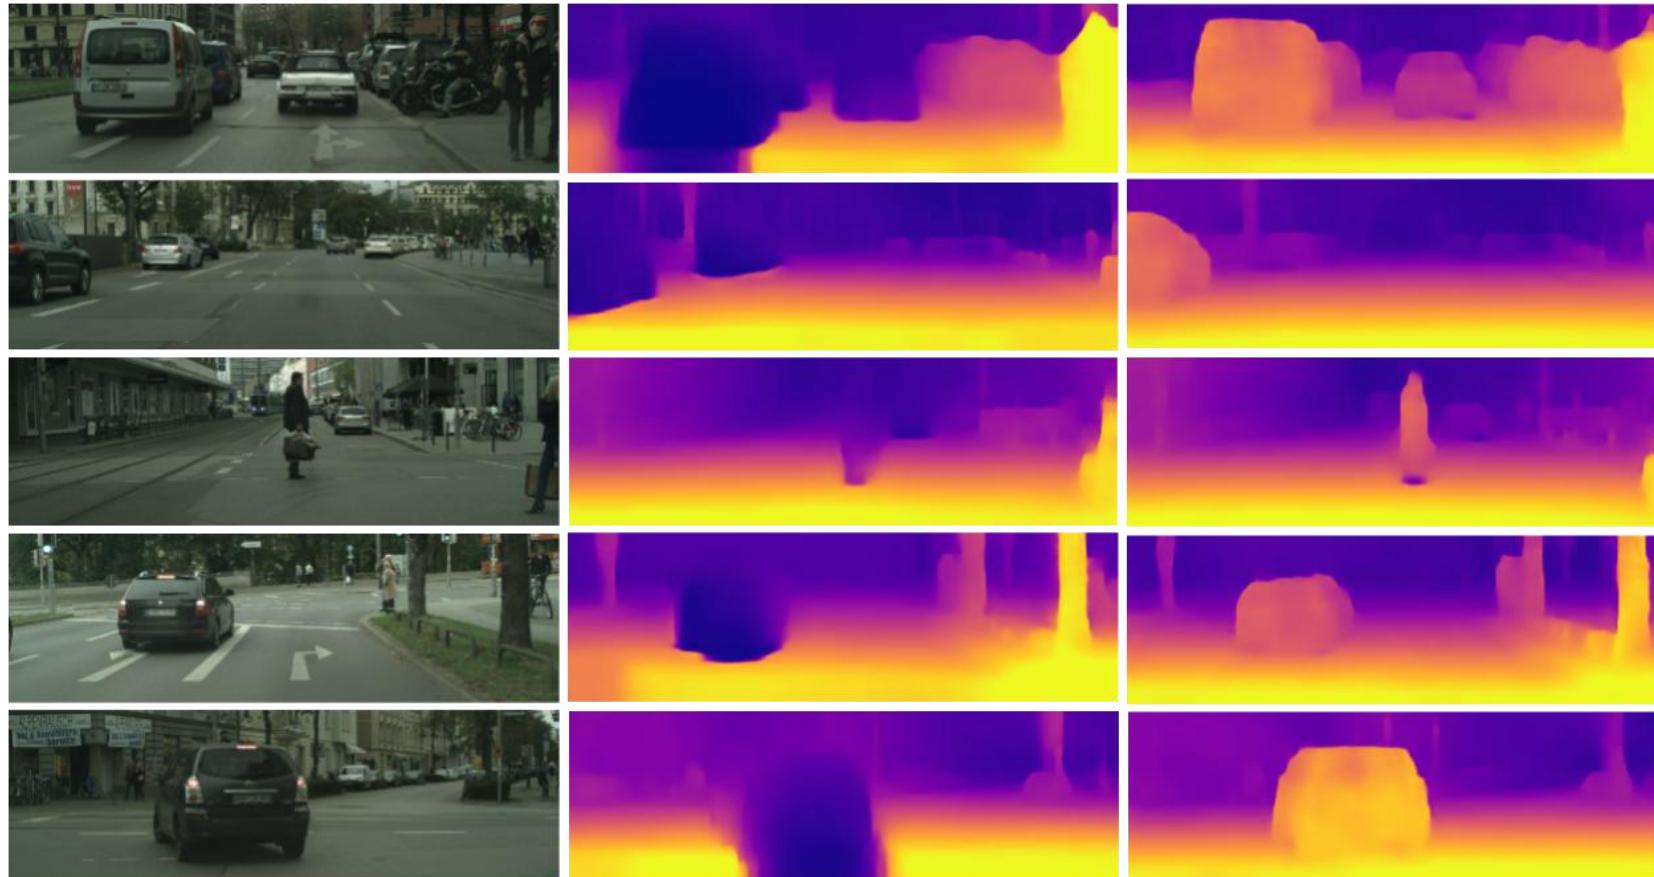
Human Pose Estimation



Bodypix



Depth Predictions



FCOS

FCOS: Fully Convolutional One-Stage Object Detection

Zhi Tian

Chunhua Shen*

Hao Chen

Tong He

The University of Adelaide, Australia

Abstract

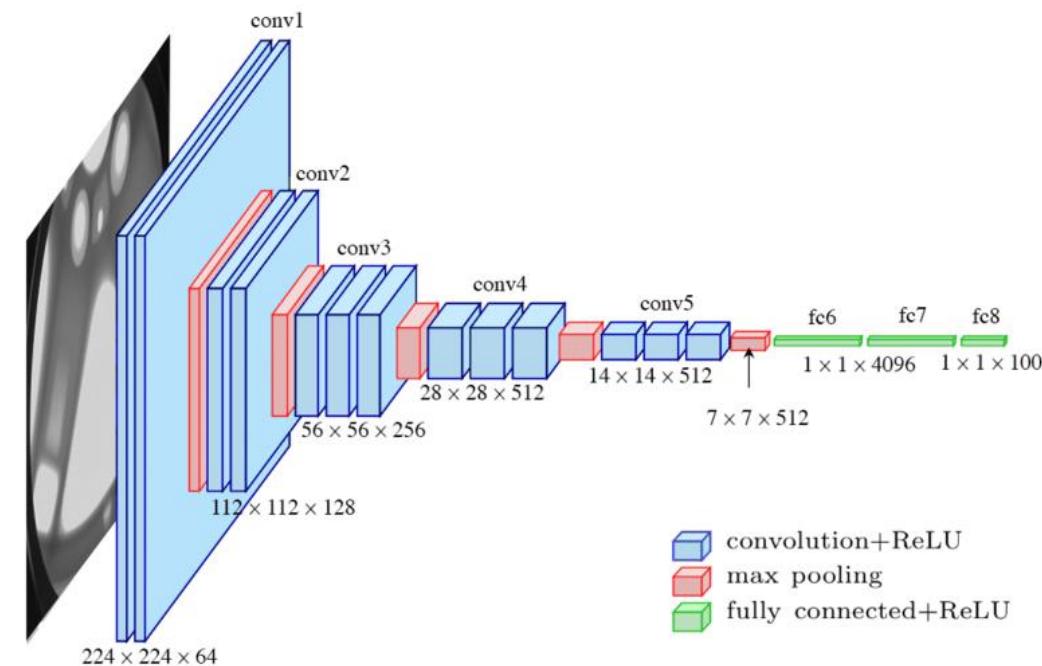
We propose a fully convolutional one-stage object detector (FCOS) to solve object detection in a per-pixel prediction fashion, analogue to semantic segmentation. Almost all state-of-the-art object detectors such as RetinaNet, SSD, YOLOv3, and Faster R-CNN rely on pre-defined anchor boxes. In contrast, our proposed detector FCOS is anchor-box free, as well as proposal free. By eliminating the predefined set of anchor boxes, FCOS completely avoids the complicated computation related to anchor boxes such as calculating overlapping during training and significantly reduces the training memory footprint. More importantly, we also avoid all hyper-parameters related to anchor boxes, which are often very sensitive to the final detection perfor-



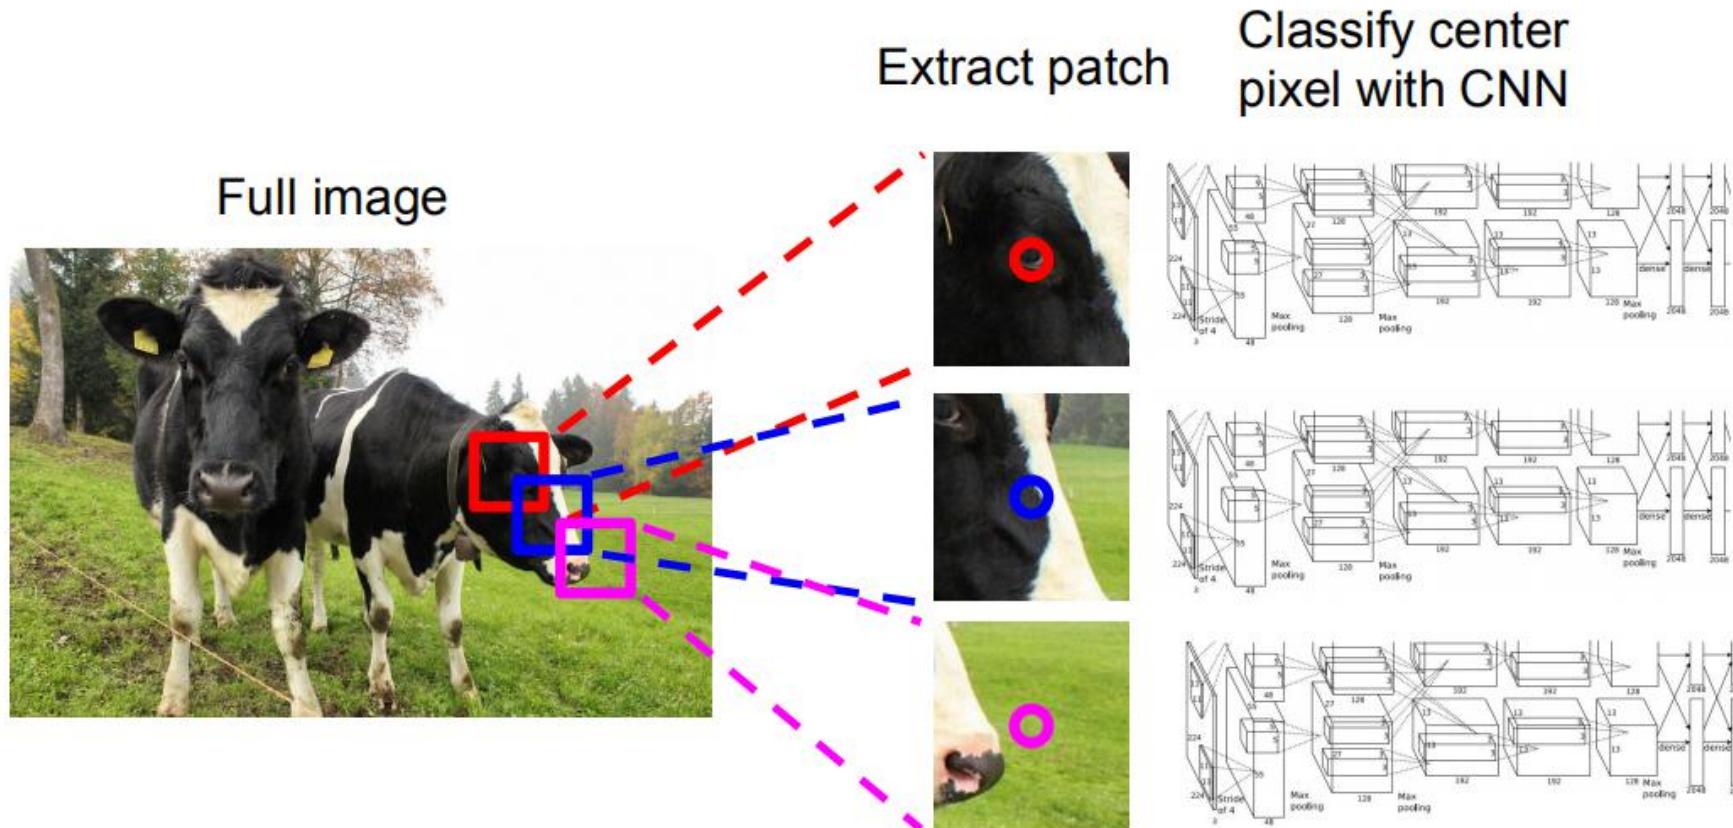
Figure 1 – As shown in the left image, FCOS works by predicting a 4D vector (l, t, r, b) encoding the location of a bounding box at each foreground pixel (supervised by ground-truth bounding box information during training). The right plot shows

Dense Predictions

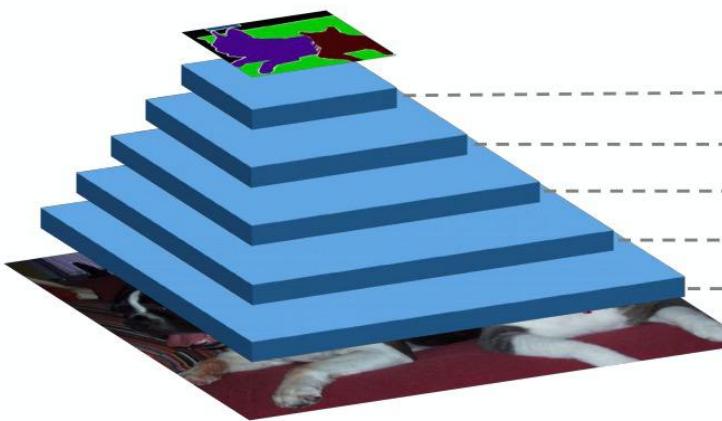
- How To do Dense Predictions from Classification?



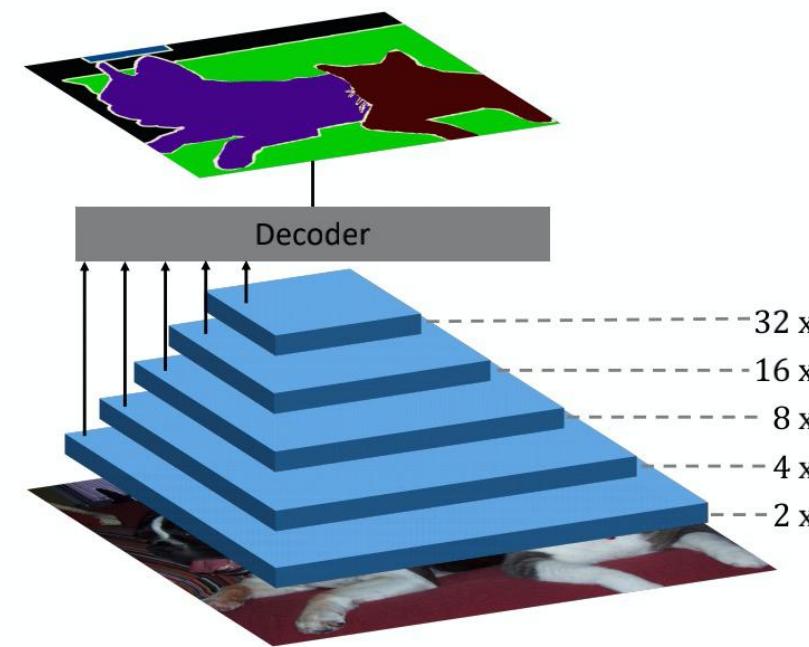
Sliding Window



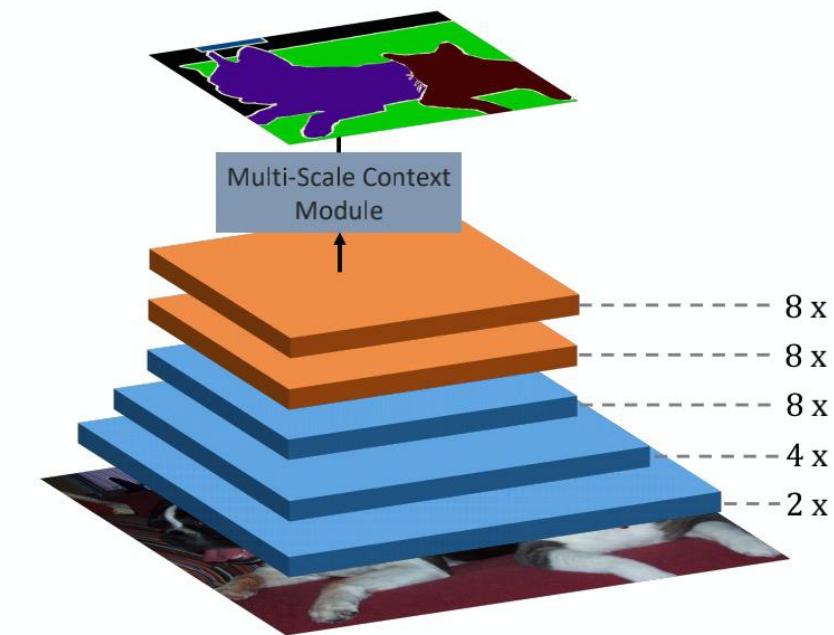
Semantic Segmentation



(a) FCN [22]



(b) EncoderDecoder



(c) DilatedFCN

HRNet

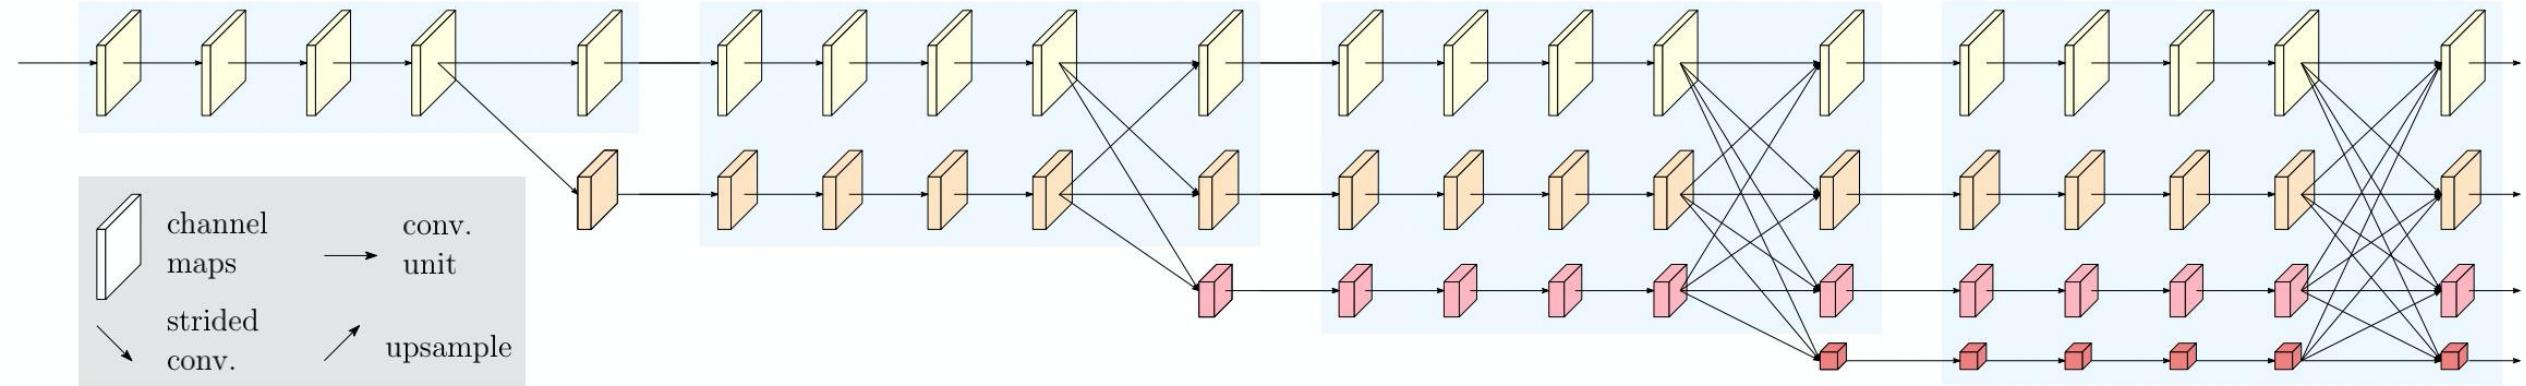
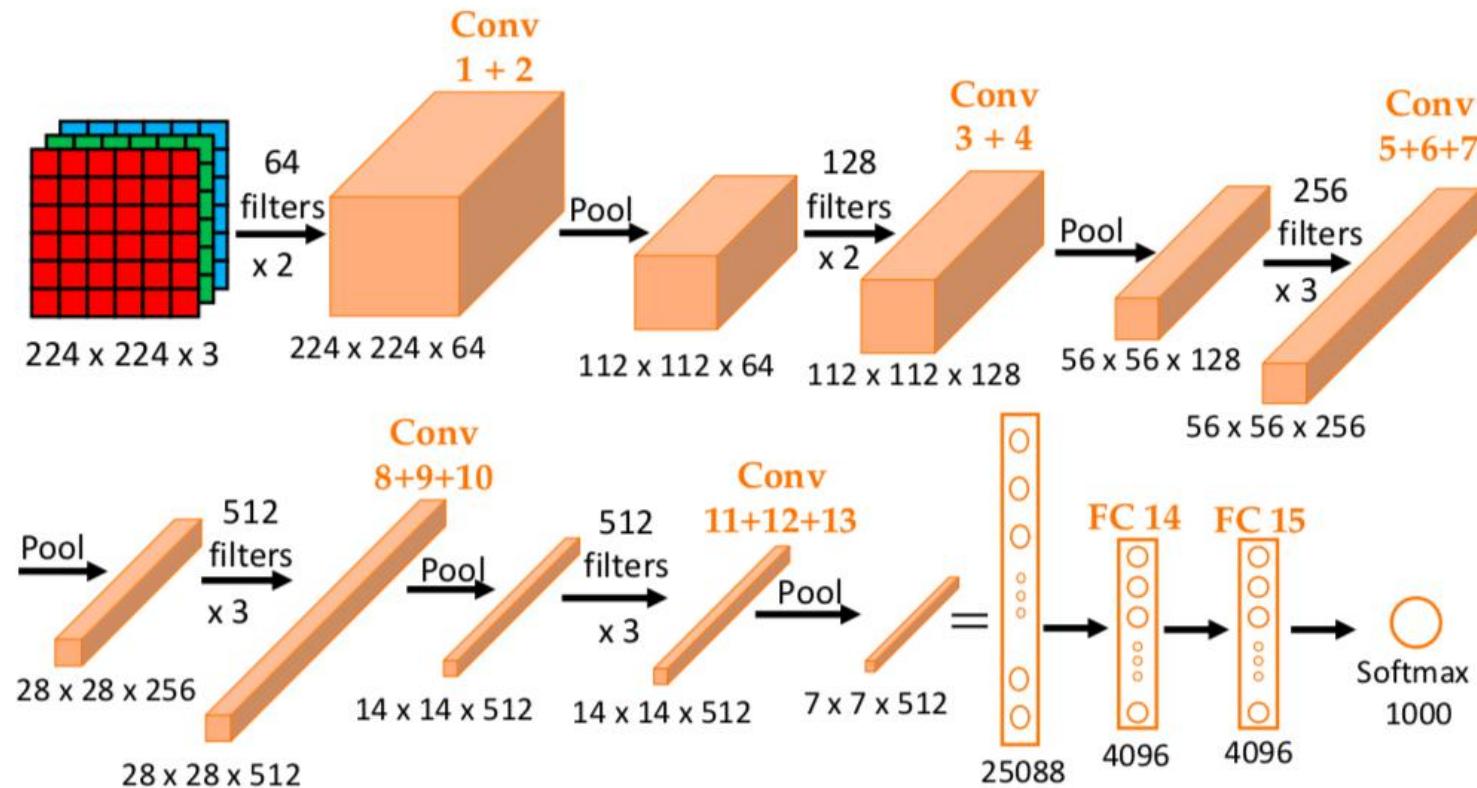


Fig. 2. An example of a high-resolution network. There are four stages. The 1st stage consists of high-resolution convolutions. The 2nd (3rd, 4th) stage repeats two-resolution (three-resolution, four-resolution) blocks. The detail is given in Section 3.

Note

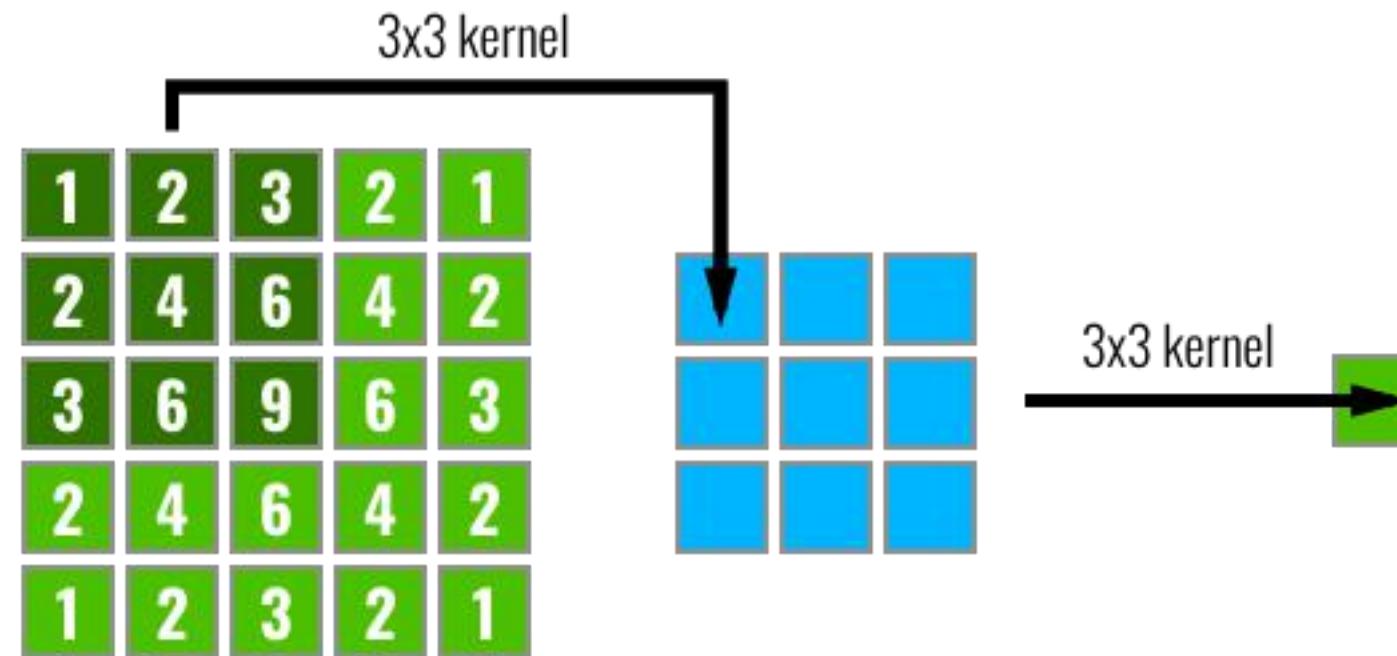
VGG



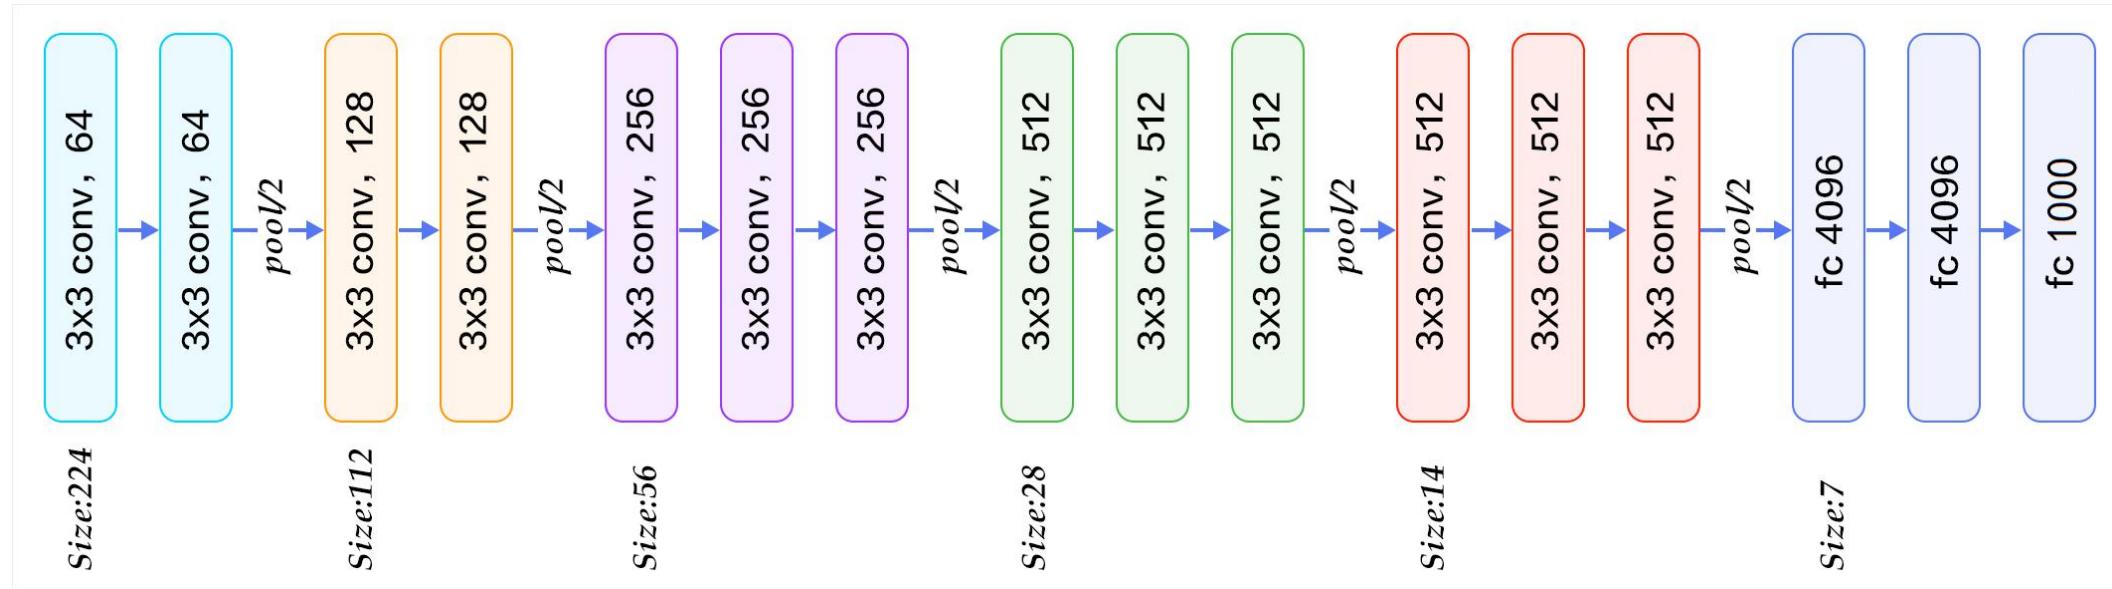
3x3 Convolution

- the smallest size to capture the notion of left/right, up/down, center

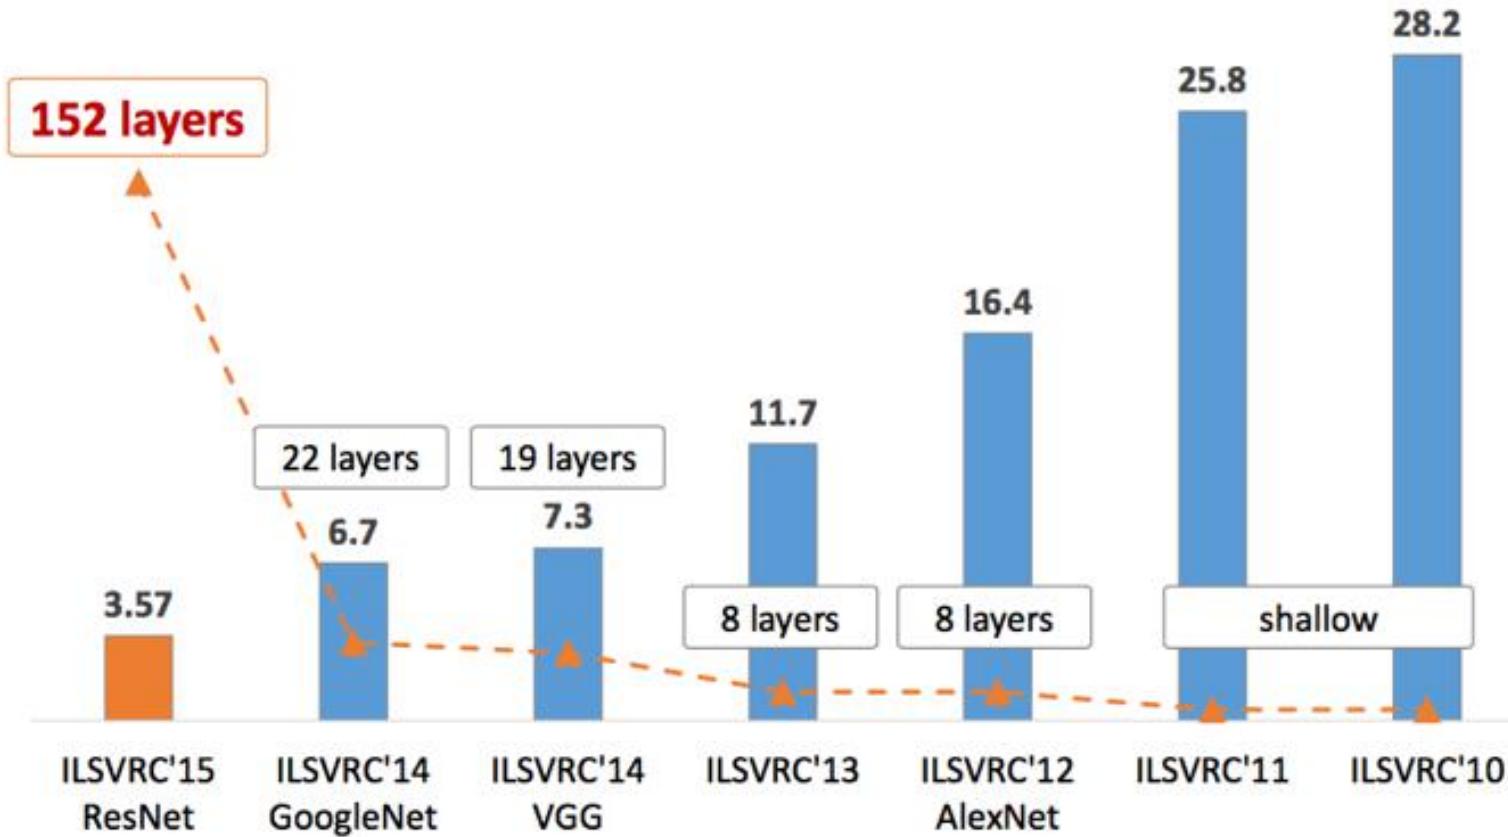
receptive field



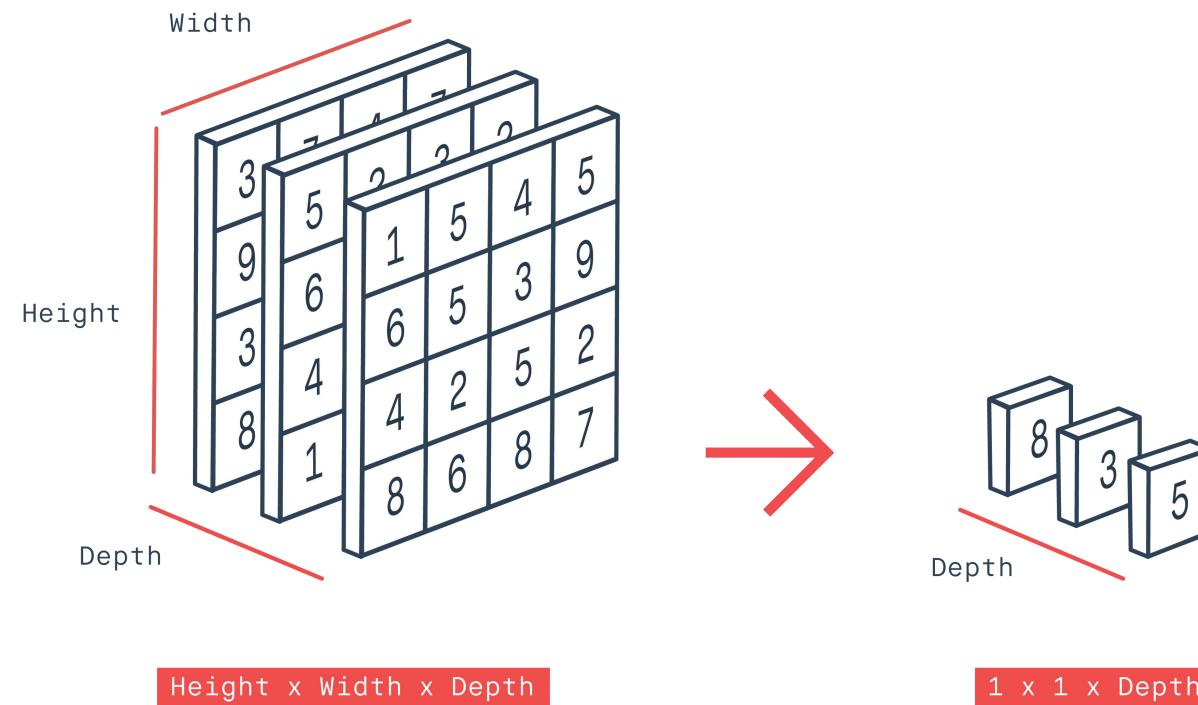
VGG



ResNet

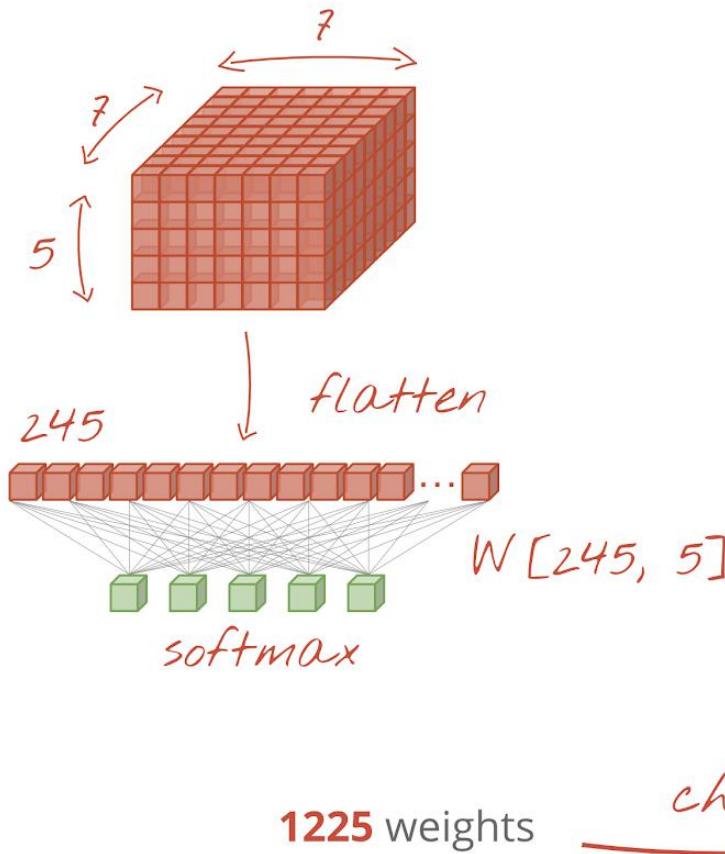


Global Average Pooling

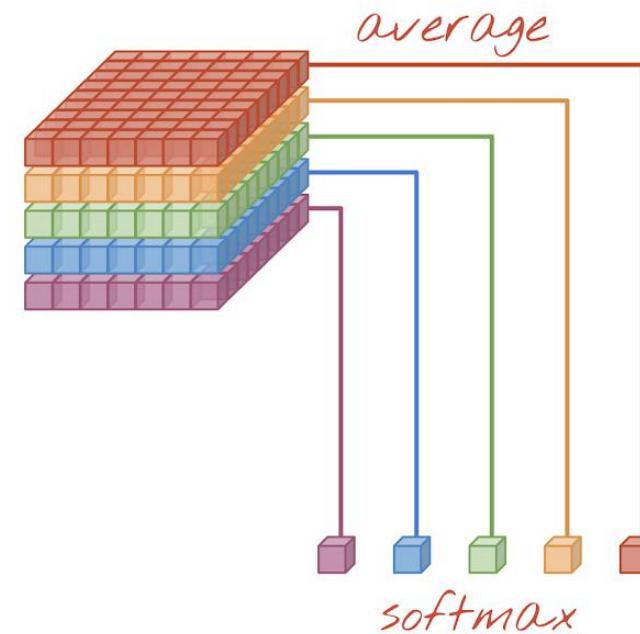


Global Average Pooling

Fully connected layer



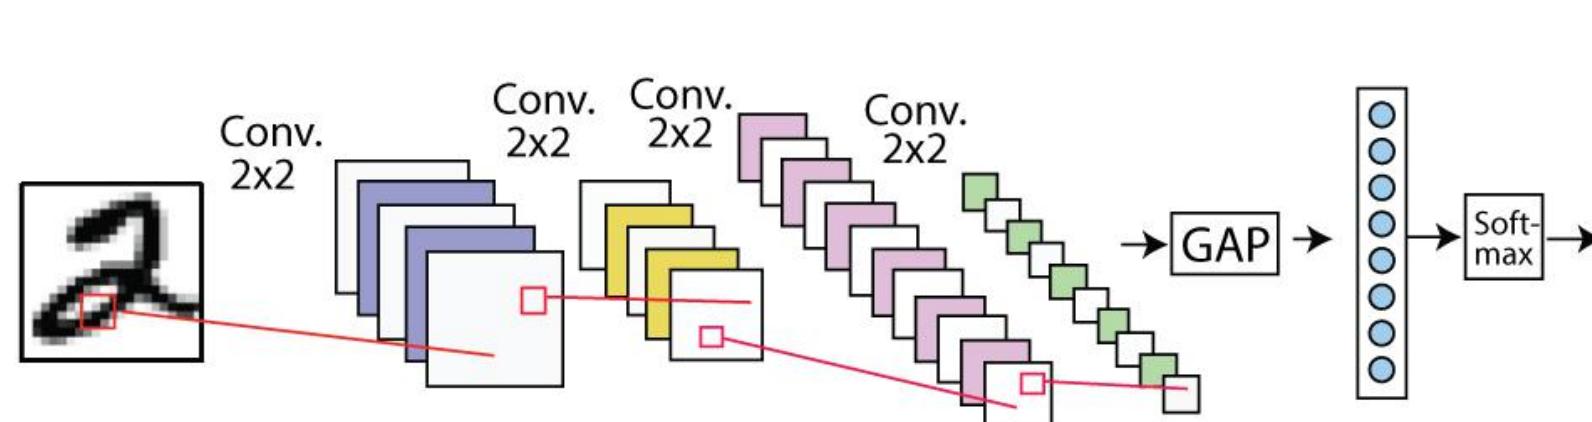
Global average pooling



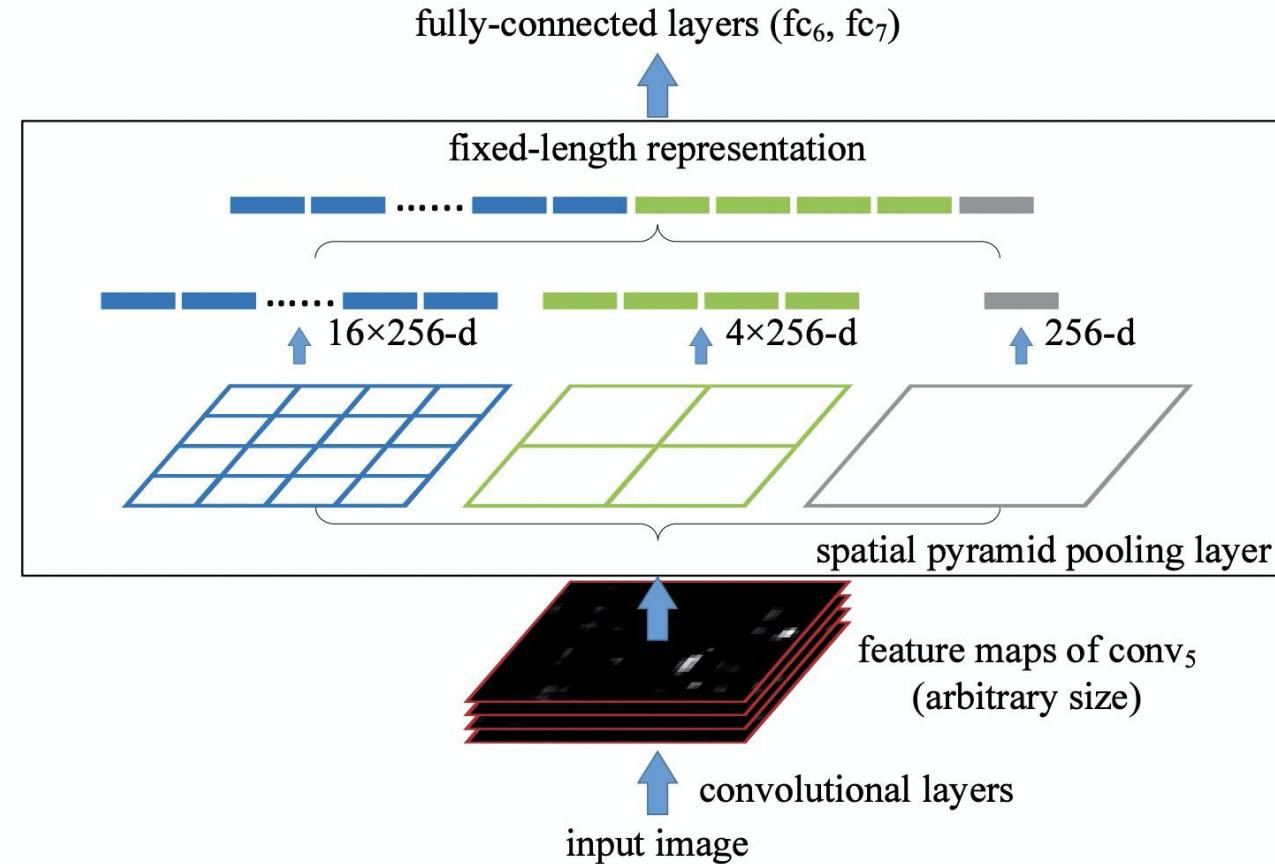
yay 
cheapskate! 

Global Average Pooling

- `torch.nn.AdaptiveAvgPool2d(output_size)`
- `tf.keras.layers.GlobalAvgPool2D`

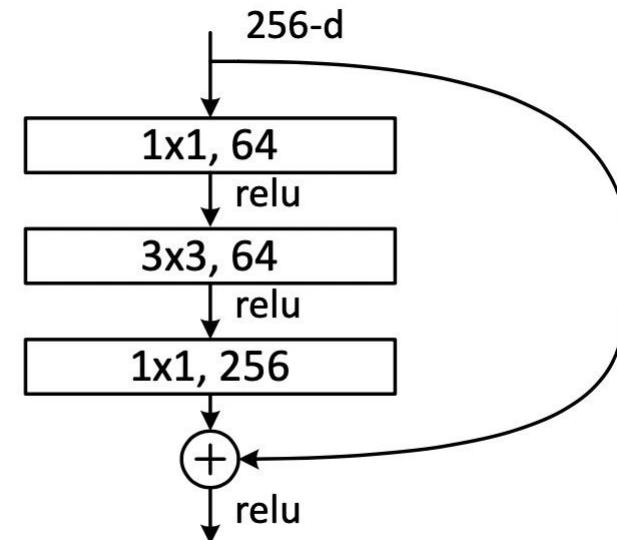
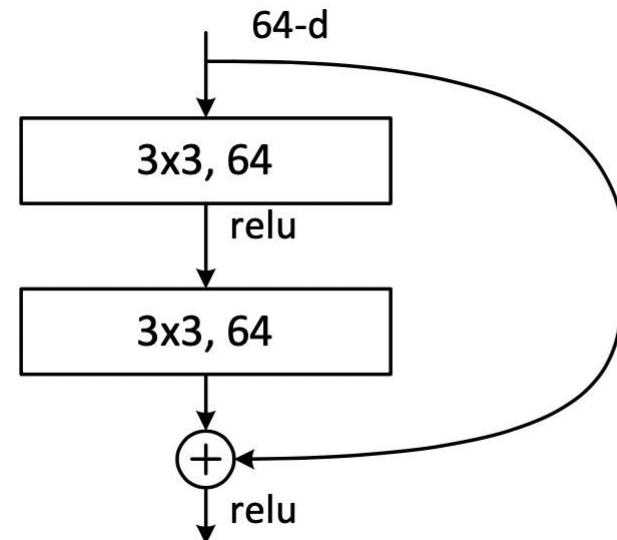


Spatial Pyramid Pooling



Notes

BottleNeck

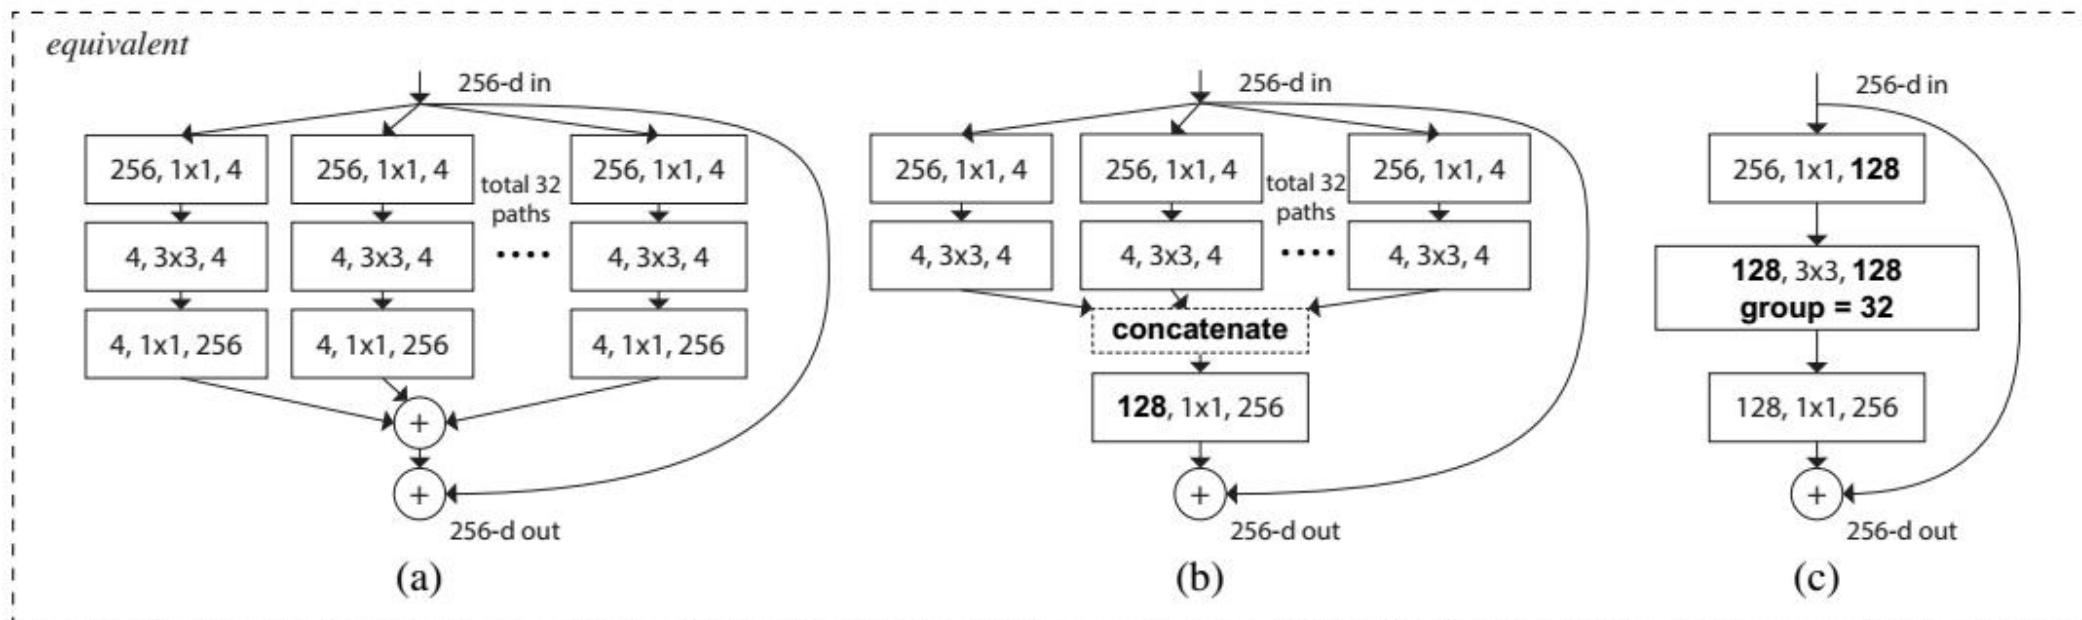


Code

ResNet新发展

- ResNeXt
- ResNeXt-Attention
- ResNeXt WSL

ResNeXt



SENet

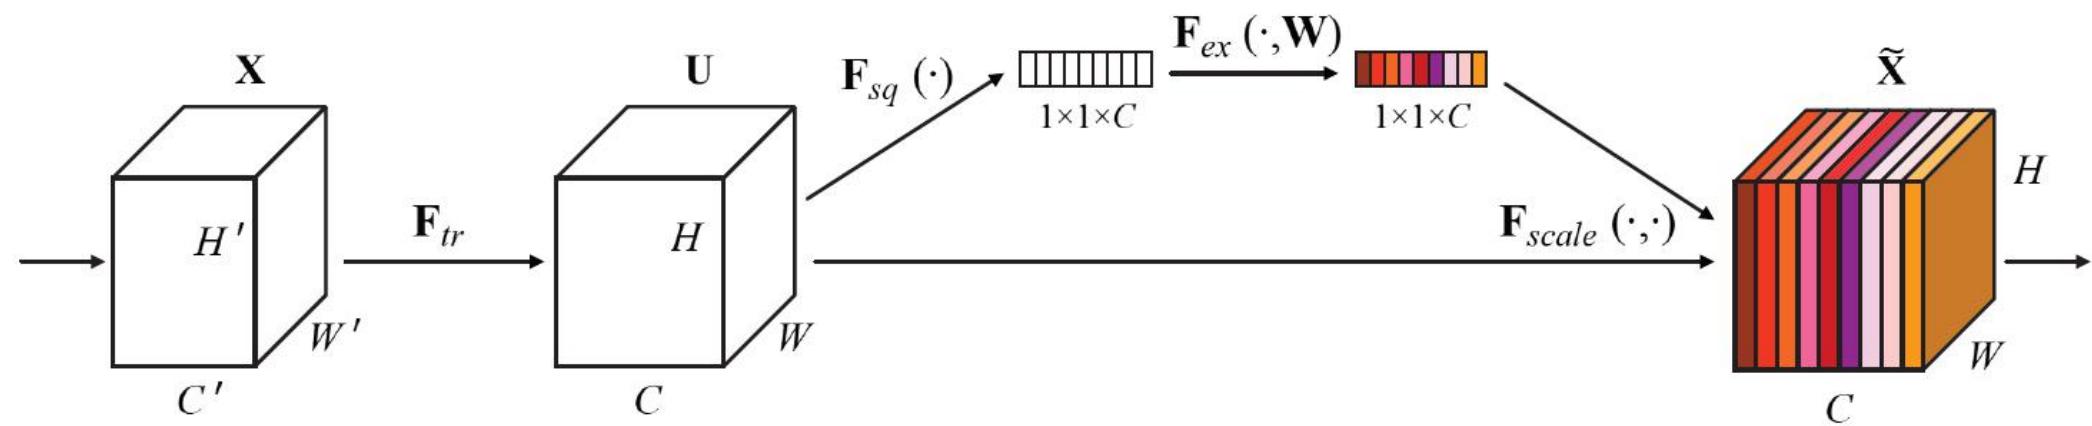
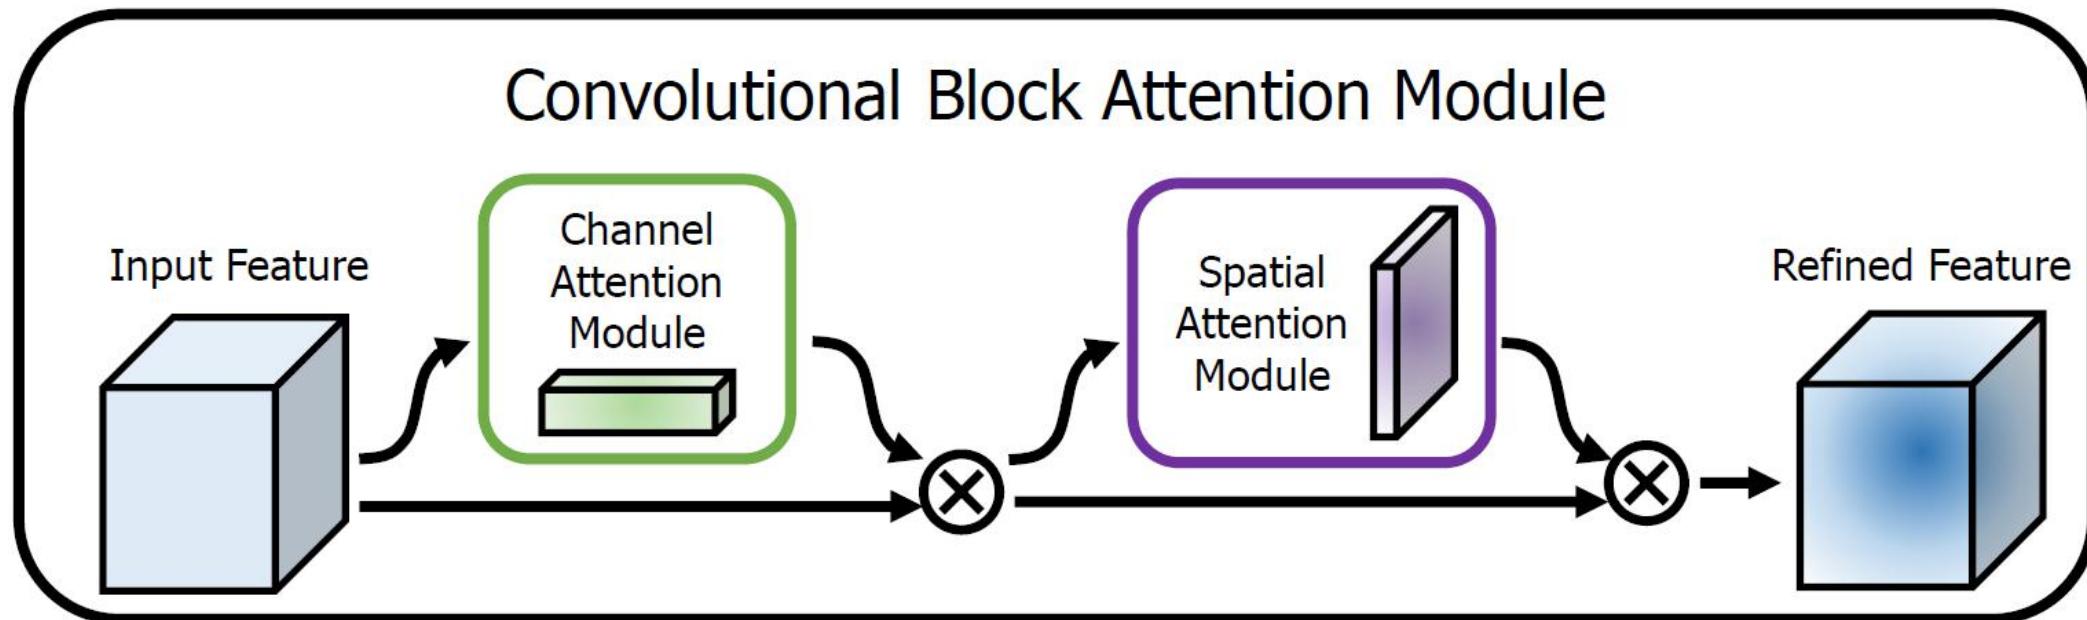


Figure 1: A Squeeze-and-Excitation block.

ResNeXt-Attention



ResNeXt WSL

PyTorch

Get Started **Ecosystem** Mobile Blog Tutorials Docs Resources GitHub Q

[View on Github](#) > [Open on Google Colab](#) >

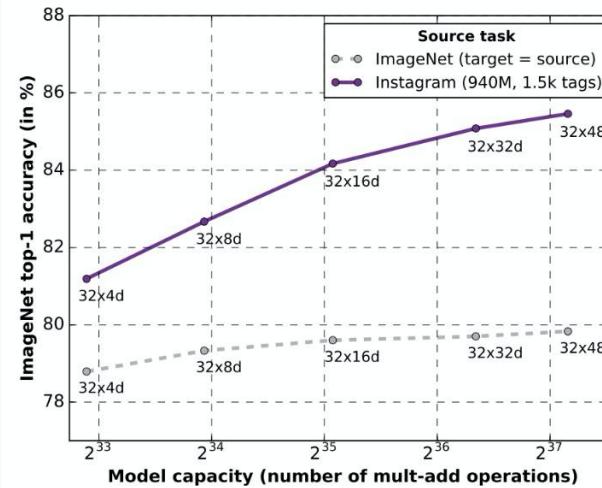


Fig. 5: Classification accuracy on val-IN-1k using ResNeXt-101 $32 \times \{4, 8, 16, 32, 48\}d$ with and without pretraining on the IG-940M-1.5k dataset.

```
import torch
model = torch.hub.load('facebookresearch/WSL-Images', 'resnext101_32x8d_wsl')
# or
# model = torch.hub.load('facebookresearch/WSL-Images', 'resnext101_32x16d_wsl')
# or
# model = torch.hub.load('facebookresearch/WSL-Images', 'resnext101_32x32d_wsl')
# or
#model = torch.hub.load('facebookresearch/WSL-Images', 'resnext101_32x48d_wsl')
model.eval()
```

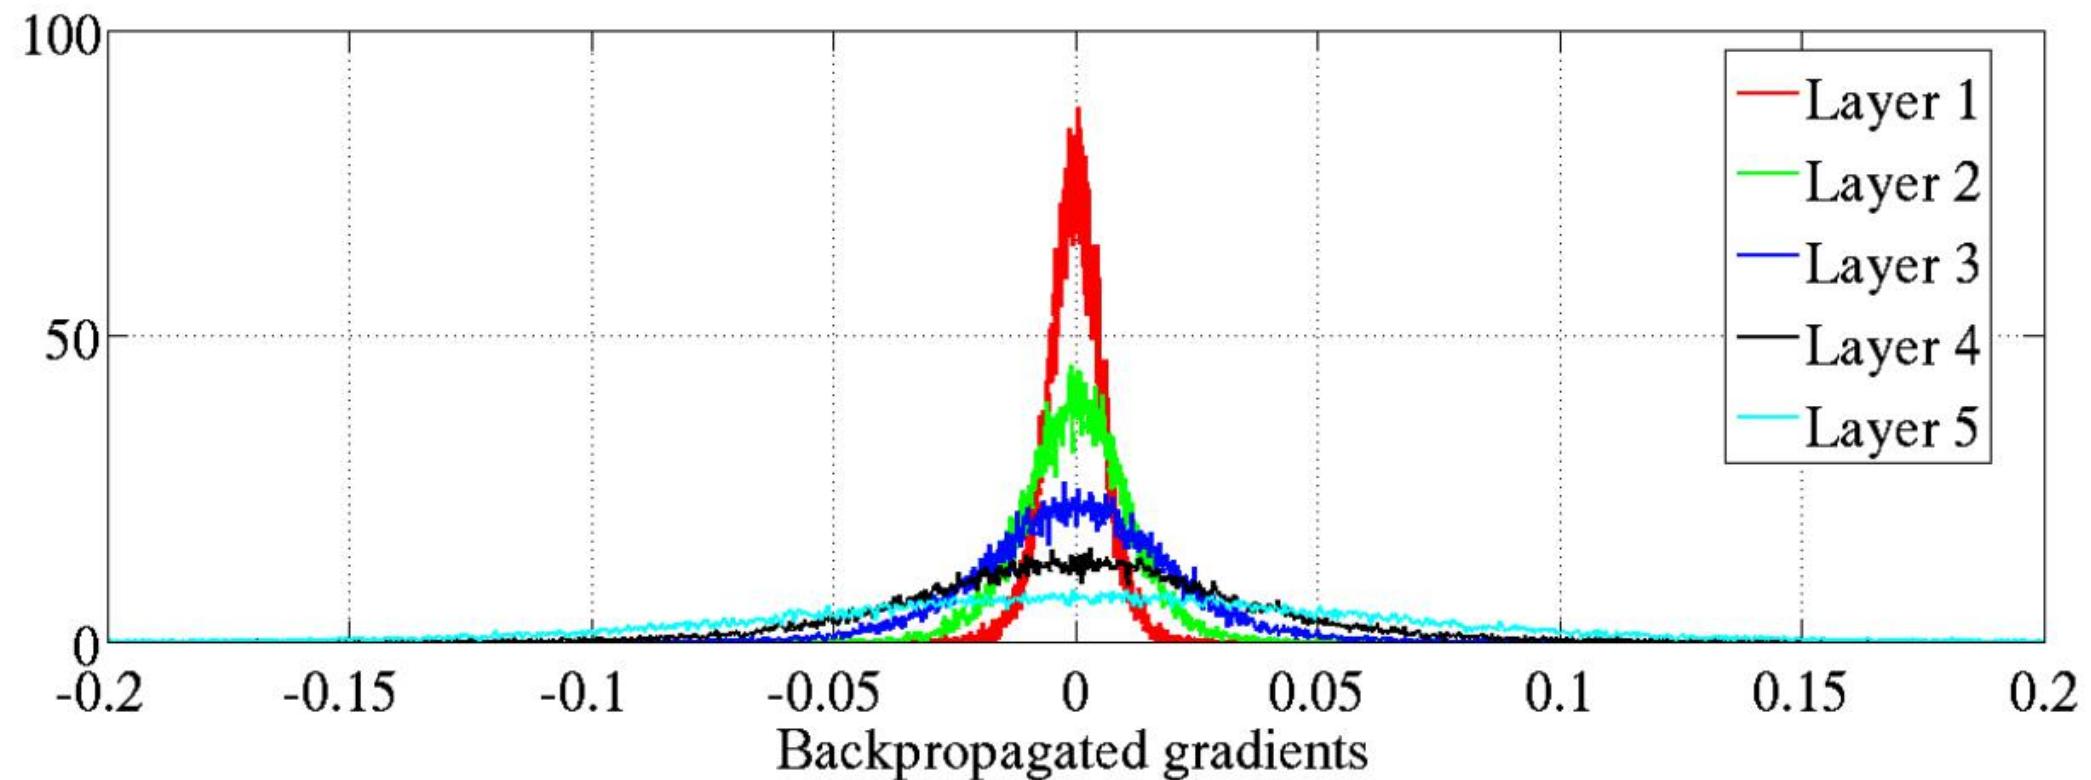
All pre-trained models expect input images normalized in the same way, i.e. mini-batches of 3-channel RGB images of shape $(3 \times H \times W)$, where H and W are expected to be at least 224. The images have to be loaded in to a range of $[0, 1]$ and then normalized using $\text{mean} = [0.485, 0.456, 0.406]$ and $\text{std} = [0.229, 0.224, 0.225]$.

Notes

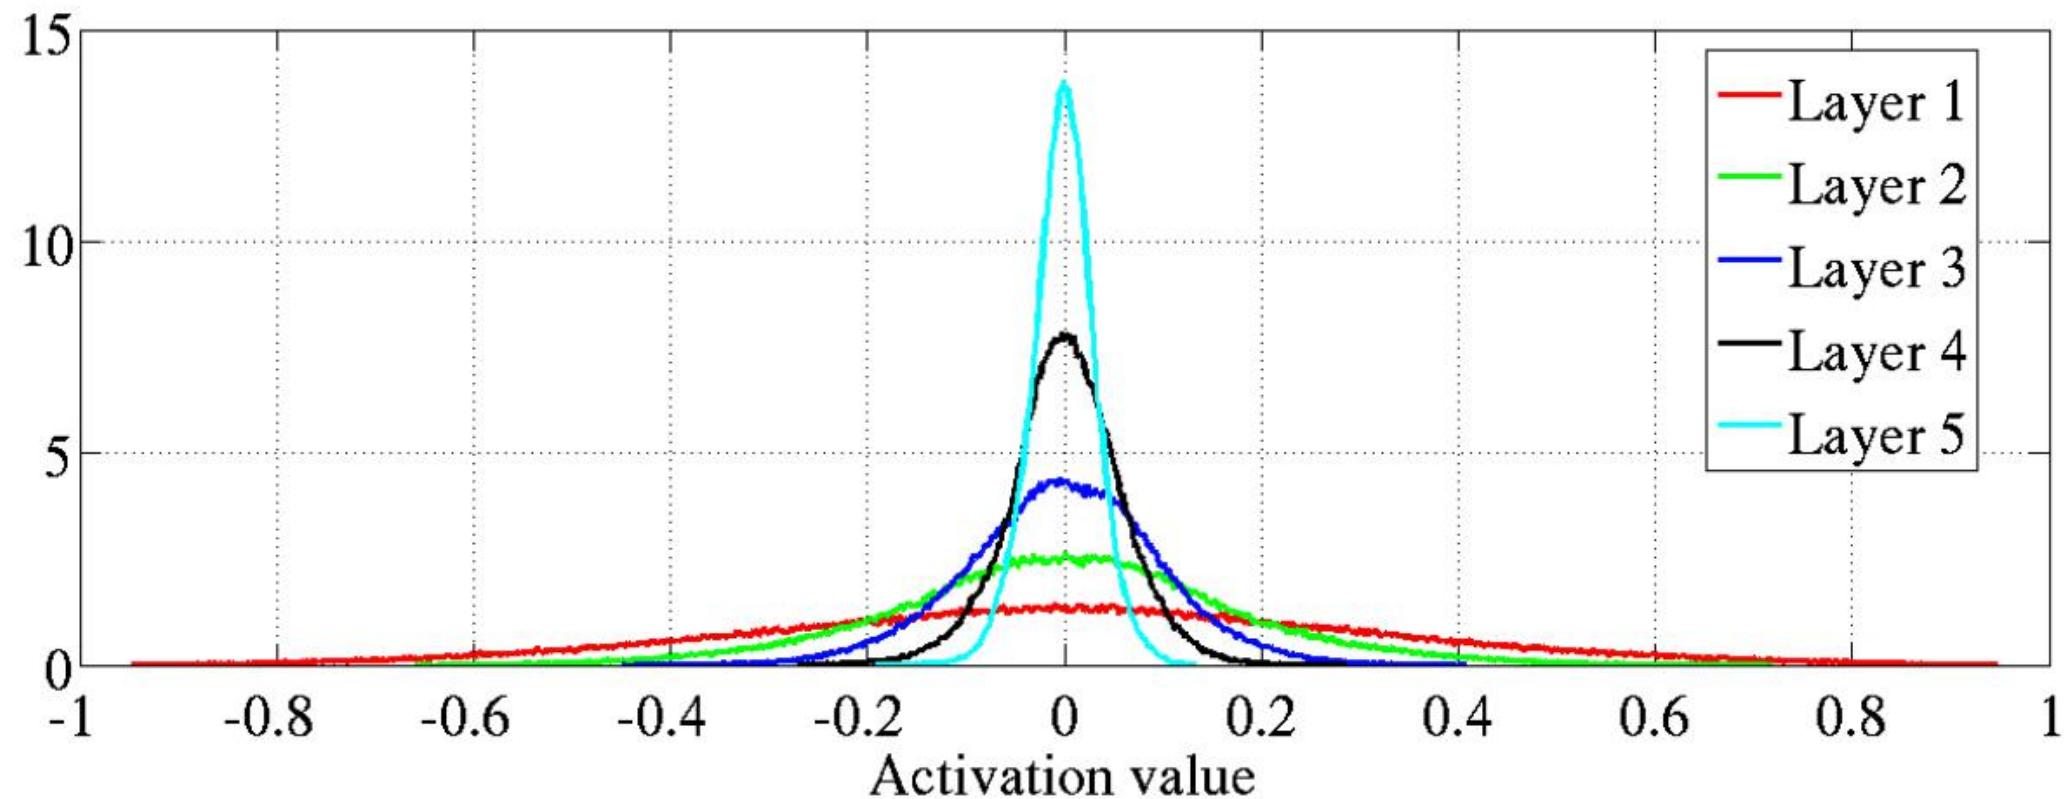
Weight Initialization

- Initialization of neural networks isn't something we think a lot about nowadays. It's all hidden behind the different Deep Learning frameworks we use, like TensorFlow or PyTorch. However, it's at the heart of why and how we can make neural networks as deep as they are today, and it was a significant bottleneck just a few years ago.

Weight Initialization



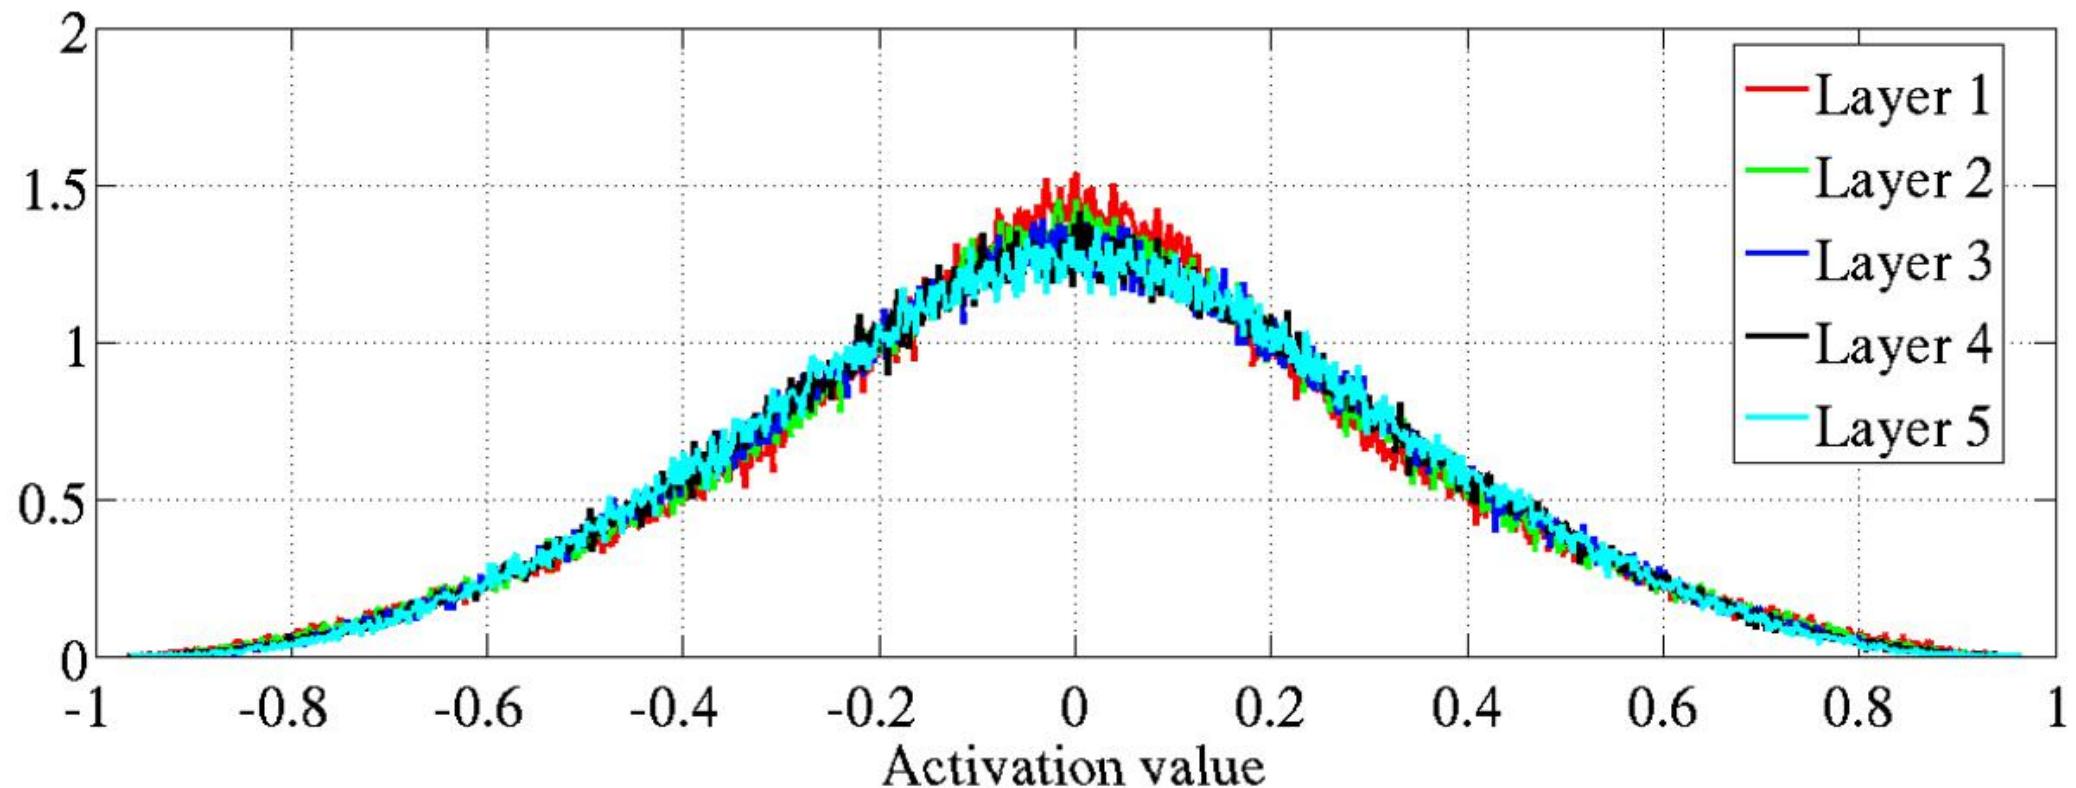
Weight Initialization



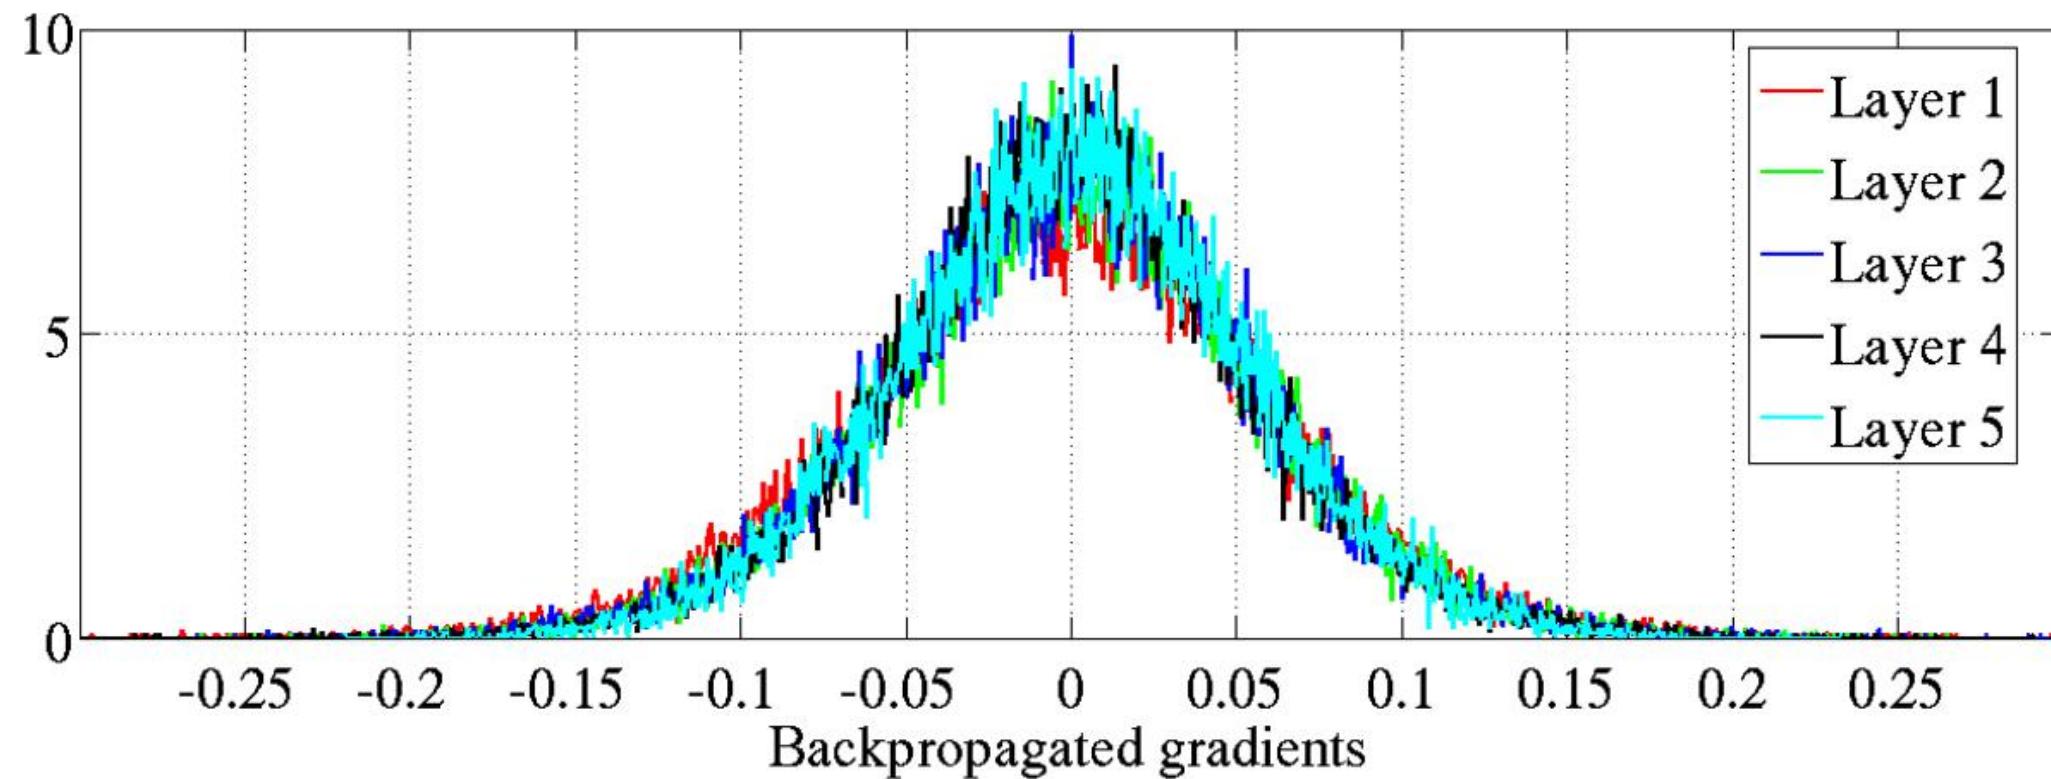
Xavier Initialization

- Understanding the difficulty of training deep feedforward neural networks
- <http://proceedings.mlr.press/v9/glorot10a.html>

Xavier Initialization



Xavier Initialization



Kaiming Initialization

- Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification
- <https://arxiv.org/abs/1502.01852>

Kaiming Initialization

$$\text{std} = \sqrt{\frac{2}{(1 + a^2) \times \text{fan_in}}}$$

Training

- Metrics
- Loss
- Optimization
- Hyperparameters

Note

课程总结

- Semantic Segmentation
- VGG
- ResNet
- Weight Initialization

重难点

- Semantic Segmentation
- VGG的实现方法
- ResNet的实现方法
- Weight Initialization的实现方法

扩展资料

- Very Deep Convolutional Networks for Large-Scale Image Recognition
- <https://arxiv.org/abs/1409.1556>
- Deep Residual Learning for Image Recognition
- <https://arxiv.org/abs/1512.03385>

扩展资料

- Convolutional Neural Networks (CNNs / ConvNets)

<https://cs231n.github.io/convolutional-networks/>

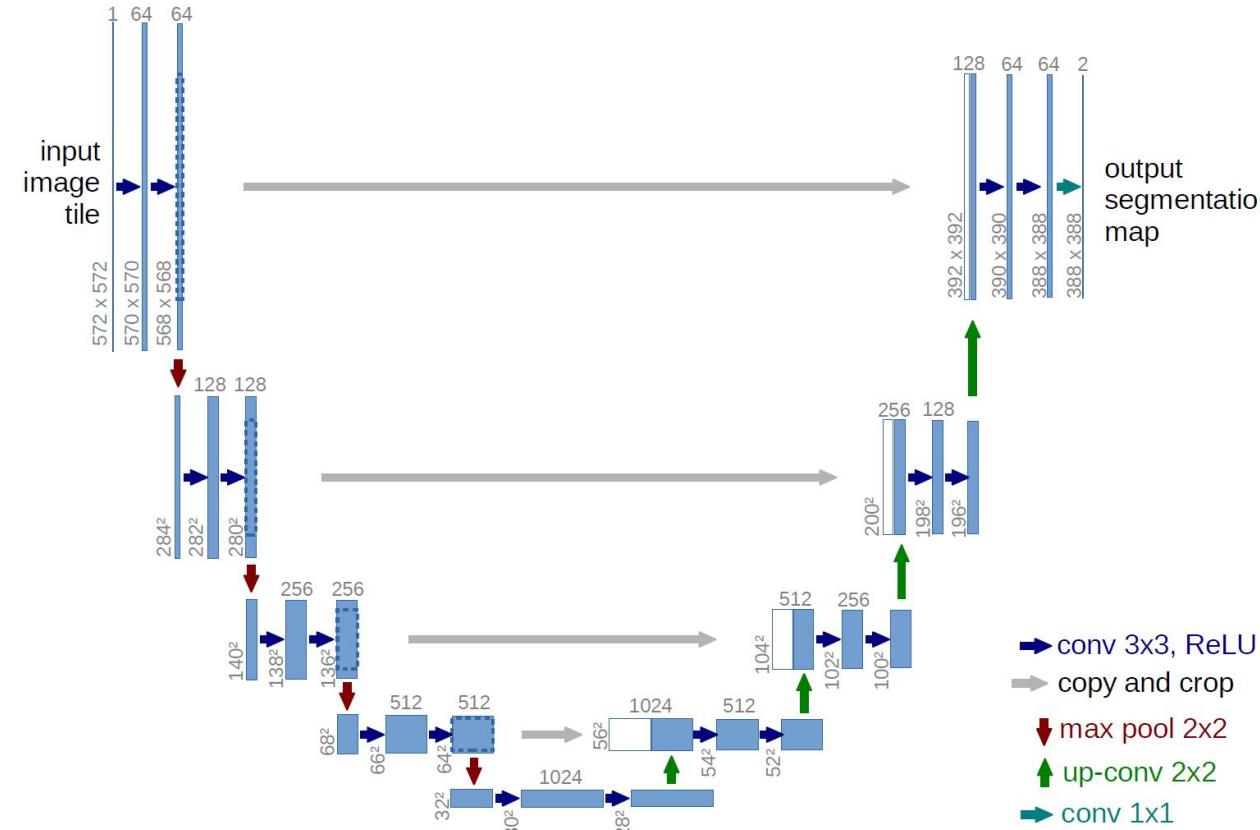
- Convolution arithmetic

https://github.com/vdumoulin/conv_arithmetic/blob/master/README.md

课程作业

- 安装开发环境
 - Framework (Tensorflow 2.0/Pytorch 1.3/PaddlePaddle 1.6)
- 实现ResNet50加载预训练参数实现图片分类

Week 2: U-Net





一所专注前沿互联网技术领域的创新实战大学